VALUATUM

20.8.2026

Advanced Micro Devices, Inc.

AMD • NASDAQ Global Select • Semiconductors • United States

COMPANY PROFILE

Advanced Micro Devices designs high-performance and adaptive semiconductors for data centers, personal computers, embedded systems and gaming. Its portfolio includes Instinct accelerators; EPYC and Ryzen processors; Pensando infrastructure products; Xilinx-derived adaptive devices; Radeon graphics processors; and semi-custom console systems-on-chip. AMD is primarily fabless, relying on external foundries, packaging providers and memory suppliers.

RECOMMENDATION	CURRENT PRICE	TARGET PRICE
SELL	**514.39 USD**	**403.80 USD**
12-month horizon	as of report date	12-month fundamental

IMPLIED UPSIDE	MARKET CAP	ENTERPRISE VALUE
−21.5%	**USD 838.8 bn**	**USD 843.9 bn**
vs. current price	shares × price	mcap + EV adjustments

P/E 2026E	EV/EBITDA 2026E	FCF YIELD 2026E	DIVIDEND YIELD 2026E	NET DEBT / EBITDA 2026E
126.4×	59.5×	-1.3%	—	0.4×

RESEARCH SNAPSHOT

The call, the math, and the three reasons why

02

A one-page summary of the recommendation and the evidence behind it; the following pages provide the detail.

RECOMMENDATION	TARGET PRICE	CURRENT PRICE	IMPLIED UPSIDE
SELL 12-month horizon	**403.80 USD** 12-month fundamental	**514.39 USD** as of report date	**−21.5%** vs. current price

MARKET CAP	ENTERPRISE VALUE	P/E · EV/EBITDA 2026E	FCF · DIV YIELD 2026E
USD 838.8 bn shares × price	**USD 843.9 bn** mcap + EV adjustments	**126.4× · 59.5×** history + peers, see p. 18–19	**-1.3% · —**

01 **The market capitalises an exceptional margin destination**

The reverse valuation requires a 66.1% EBIT margin in 2027, versus 33.4% in the followed model. AMD needs more than revenue growth: it must turn challenger hardware into platform economics.

66.1% implied 2027 EBIT margin

02 **Accelerator evidence is strong but not yet ecosystem-complete**

MI355X benchmarks, customer deployments and the Meta commitment validate demand. However, lower cited multi-GPU scaling efficiency and CUDA's continuing ecosystem advantage leave terminal share and margin outcomes uncertain.

5–10% external share range

03 **Diversification supports value but cannot carry the premium**

EPYC, Ryzen and adaptive computing are valuable franchises with improving share or recovery potential. Even generous peer-bracket multiples for these businesses leave accelerators responsible for most of the current enterprise value.

71.7% of EV in accelerators

THESIS BREAKER

The SELL thesis fails if scaled MI400 deployments generate recurring accelerator EBITDA above USD 15bn and group EBIT margins remain sustainably above 35%.

SHARE PRICE DEVELOPMENT

Price trajectory and valuation anchors

03

Weekly closing prices over 36 months with the 12-month target price, alongside key market data.

Share price — last 36 months

USD per share, with 12-month target line

CURRENT PRICE **514.39 USD** 20 August 2026	TARGET PRICE **403.80 USD** 12-month	IMPLIED UPSIDE **−21.5%** vs. current	52-WEEK HIGH **584.73 USD** past 12m
52-WEEK LOW **149.22 USD** past 12m	1-YEAR CHANGE **+182.3%** price only	3-YEAR CHANGE **+342.3%** price only	REPORT DATE **20 August 2026**

READING THE MOVE

The supplied annual controls show market capitalisation rising from USD 349.1bn at 2025 year-end to USD 838.8bn for the 2026 control period. The evidence points to three fundamental drivers: rapid Data Center growth, visible Instinct qualifications and multi-gigawatt commitments, and improving EPYC share. It does not establish the exact drivers of day-to-day price moves. Our target assumes continued execution, but not the full margin outcome embedded in the share price.

VALUE BREAKDOWN · PART 1

04

Enterprise-value allocation by business division

Where the market value sits across current businesses and long-term strategic options.

Enterprise value by business division

USD millions · current analytical allocation

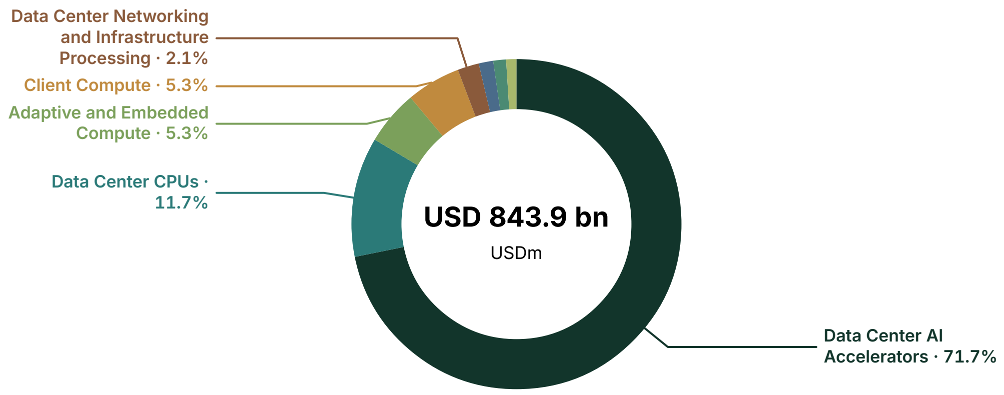

VALUE AND EARNINGS BY BUSINESS DIVISION

Analytical enterprise-value allocation against reported economics. "Reported" rows match a verified company-reported segment; "Modelled allocation" rows are model-side splits. Business divisions may differ from company-reported accounting segments.

	EV	EV %	REVENUE	REV %	REPORTED EBIT	REPORTED EBIT % OF TOTAL	BASIS
Data Center AI Accelerators	605,000	71.7%	9,000	26.0%	1,200	32.5%	Modelled allocation
Data Center CPUs	99,000	11.7%	7,000	20.2%	1,000	27.1%	Modelled allocation
Adaptive and Embedded Compute	45,000	5.3%	3,500	10.1%	700	18.9%	Modelled allocation
Client Compute	44,800	5.3%	10,600	30.6%	900	24.4%	Modelled allocation
Data Center Networking and Infrastructure Processing	17,600	2.1%	1,200	3.5%	100	2.7%	Modelled allocation
Radeon Discrete Graphics	11,961	1.4%	904	2.6%	-100	-2.7%	Modelled allocation
Semi-Custom Gaming SoCs	10,500	1.2%	3,000	8.7%	194	5.3%	Modelled allocation
AI Infrastructure Systems	10,000	1.2%	435	1.3%	-300	-8.1%	Modelled

VALUE BREAKDOWN · PART 2

05

Where value sits versus where money is earned

Each business division's share of enterprise value against its share of group revenue and profit. A division holding value it does not yet earn is where the valuation rests on the future; the full figures are in the allocation table on the previous page.

Share of enterprise value vs revenue and profit

Per cent of group total · modelled allocation

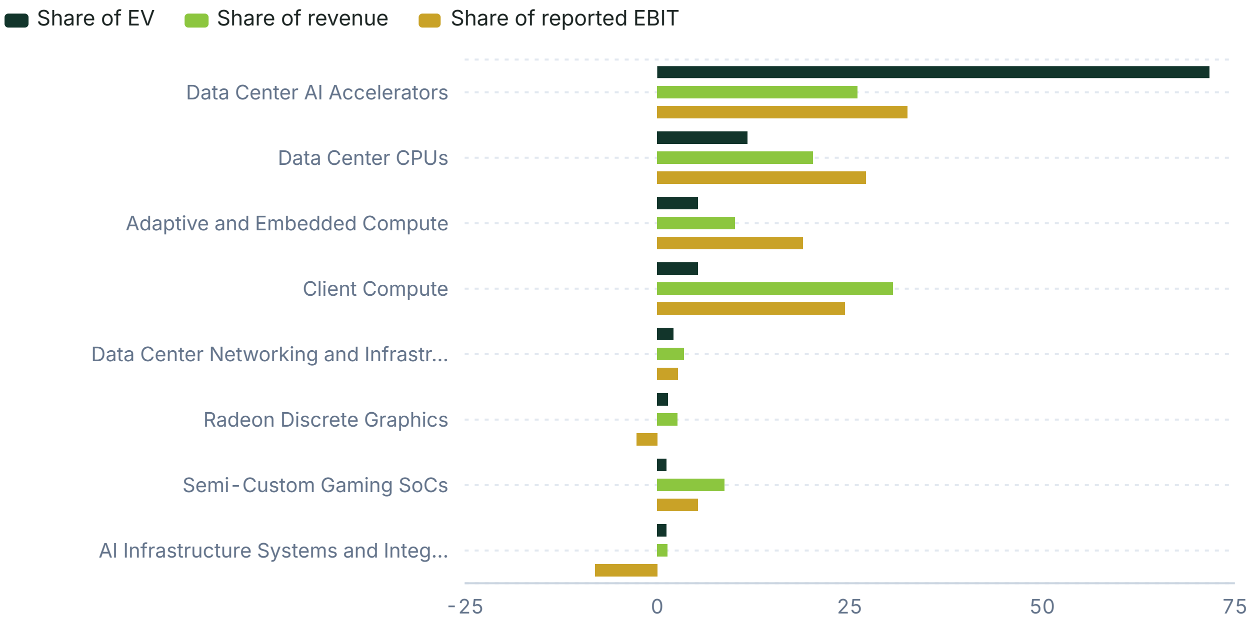

INTERPRETATION

The control enterprise value is concentrated in Data Center AI Accelerators at USD 605.0bn, or 71.7%. Data Center CPUs account for 11.7%, while Client and Adaptive/Embedded each represent 5.3%. This is a current-market allocation, not management disclosure. Non-accelerator franchises offer meaningful downside support, but cannot justify the current valuation if accelerator margins or deployment scale disappoint.

RECOMMENDATION

Why we land on the call we do

06

How the assembled evidence adds up to the call: what the price assumes, what the clues say, and what would change our view. The decision summary and the three headline reasons are on the snapshot page.

SELL. The 12-month target price is USD 403.8, implying 21.5% downside from the current price of USD 514.39.

At the current price, most of AMD's enterprise value lies in Data Center AI Accelerators. The valuation effectively requires AMD to approach NVIDIA-like operating economics. The reverse valuation requires 2027 EBIT of USD 52.5bn on USD 79.4bn of revenue, a 66.1% EBIT margin, versus USD 26.5bn of EBIT in the followed model. This gap matters more than quarterly revenue beats of several hundred million dollars.

The evidence is constructive. MI355X achieved 119% of B200 interactive performance in one MLPerf Llama 2 70B comparison, Meta committed to a multi-year deployment of up to 6GW, and Q2 Data Center revenue reached USD 6.7bn. However, external accelerator-share estimates range from 5–7% to approximately 10%, cited multi-GPU scaling tests remain weaker than NVIDIA's, and export rules limit the highest-performance products.

The strongest argument against SELL is that a second merchant accelerator platform could capture disproportionate value from a USD 653.4bn annual data-center-systems spending pool, especially as customers seek supply diversity. This could justify the current price if AMD sustains roughly USD 15bn of 2027 accelerator EBITDA and proves that rack deployment, ROCm portability and supply scale reinforce one another. It does not yet support the reverse valuation's 66% EBIT margin.

A sustained move above the target would require repeat accelerator orders beyond initial qualifications, ROCm production performance close to CUDA across multi-node workloads, positive free cash flow despite capacity commitments, and a clear path to at least 35% group EBIT margins. The downside case strengthens if MI400 deployments slip, hyperscalers shift spending to custom ASICs, EPYC share stalls against Arm, or inventories and working capital continue to absorb higher earnings.

CORE ANALYSIS · EXECUTIVE MAP

07 How the company creates economic value

The framework underlying the business division deep dives that follow.

Instinct drives enterprise value. The key question is whether modelled 2027 accelerator EBITDA can reach USD 15.0bn. Hardware competitiveness and customer commitments support the revenue case, but ROCm, rack integration and pricing must turn that revenue into durable platform margins. EPYC is the second-largest business division, with USD 4.5bn of modelled 2027 EBITDA. It is supported by 46.2% x86 server revenue share but challenged by Arm's 17.7% data-center unit share.

Client Compute and Adaptive/Embedded provide diversified cash earnings. The key indicators are 31.4% client x86 revenue share and Q2 Embedded growth of 19%. Networking and systems integration are attachment businesses, with value dependent on standard content in AMD racks. Semi-custom has secured important next-generation designs but must bridge a mature console cycle. Radeon must turn an 8% AIB share into positive standalone economics.

The authoritative group model forecasts revenue rising from USD 49.7bn in 2026 to USD 108.9bn in 2028. FCFF remains negative through 2027 because capacity, working capital and investment absorb cash. The decisive group metric is therefore not revenue growth alone, but conversion of USD 26.6bn of 2027 EBITDA into repeatable free cash flow.

ANALYTICAL ANCHOR

At USD 514.39, AMD is valued less as a portfolio of challenger franchises and more as an accelerator platform with a proven margin outcome.

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.1 Data Center AI Accelerators

Business division deep dive — Profit engine.

THE OPPORTUNITY IS THE DATA-CENTER BUDGET BEING REALLOCATED, NOT A NARROWLY DEFINED ACCELERATOR MARKET

Gartner forecasts annual worldwide spending on data-center systems of USD 653.4bn in 2026. This is the broad infrastructure pool against which accelerator systems compete. A separate estimate puts the 2026 AI-accelerator market at USD 240bn, while another labelled-market forecast reaches only USD 63.2bn by 2030. These definitions cannot be used interchangeably. The valuation uses broader displacement logic because accelerator adoption changes spending on servers, networking, memory, power and software. It does not assume AMD captures rack revenue simply because its GPU drives that spending. The report's modelled USD 34.0bn of 2027 accelerator revenue equals 5.2% of the broad 2026 systems-spending pool and 14.2% of the narrower USD 240bn accelerator estimate. This is an ambitious but visible capture requirement, not a claim that AMD already has that share.

The market is expanding quickly enough for AMD to grow without immediately displacing every NVIDIA installation. Gartner estimates 2026 data-center electricity consumption at 565TWh, including 175TWh for AI-optimised servers. It expects AI-server consumption to reach 258TWh in 2027. Combined 2026 capital expenditure for Amazon, Microsoft, Alphabet and Meta is estimated at USD 725–760bn. As the sources differ, the analysis uses the lower USD 725bn boundary. Not all capex buys merchant GPUs, and an external estimate assigns custom ASICs 25% of accelerator revenue. AMD must win within merchant accelerators while defending the category against TPU, Trainium, MTIA and other internally designed silicon.

THE CURRENT VALUATION REQUIRES AMD TO MOVE FROM SECOND SOURCE TO DURABLE PLATFORM

The report models accelerator revenue rising from approximately USD 9.0bn in FY2025 to USD 34.0bn in 2027. EBITDA reaches USD 15.0bn, a 44.1% margin. The FY2025 starting point is not company-reported product revenue. It is the report's allocation within the officially reported USD 16.635bn Data Center segment, informed by accelerator-revenue commentary and the need to value EPYC, networking and systems integration separately. External estimates differ materially. Silicon Analysts places AMD at 5–7% of the 2026 accelerator market, while IDC places it near 10%. The base case weights the lower range because the 10% estimate is hard to reconcile with NVIDIA's reported scale and AMD does not disclose product revenue.

DATA CENTER AI ACCELERATORS FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 9,000m
MODEL ESTIMATE REPORTED EBIT	USD 1,200m
FY2025 MODELLED REVENUE	$9.0bn
2027E MODELLED EBITDA	$15.0bn
CURRENT-MARKET SOTP MULTIPLE	40.33x 2027E EBITDA
EXTERNAL 2026 SHARE ESTIMATES	5–7% to ~10%
MI355X MEMORY CAPACITY	288GB HBM3E
MI300X CLOUD RENTAL RANGE	$1.71–$6.00/GPU-hour
MODELLED 2027 ACCELERATOR REVENUE	$34.0bn
META DEPLOYMENT COMMITMENT	Up to 6GW

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.1 Data Center AI Accelerators (continued)

Business division deep dive — Profit engine.

At USD 34.0bn of revenue, a blended accelerator ASP of USD 25,000 implies approximately 1.36m accelerator-equivalent units. This is the report's estimate, not a shipment disclosure. The physical mix will include modules and rack configurations, so a simple card count is imperfect. The cited MI300X price range of USD 10,000–20,000 and the MI400 forecast ASP of USD 30,926 imply a plausible 2027 range of roughly 1.1m to 2.3m units for the same revenue. The wide range matters. The high-unit outcome puts more pressure on HBM and advanced packaging, while the high-ASP outcome requires AMD to capture more system value and show that customers will pay for software and deployment quality, not memory capacity alone.

HARDWARE IS CREDIBLE, BUT PERFORMANCE ACROSS WORKLOADS REMAINS THE TEST

MI355X provides 288GB of HBM3E and 8TB/s peak bandwidth, compared with 192GB on the cited standard NVIDIA B200 configuration. In MLPerf Inference 6.0, AMD's MI355X platform reached 119% of B200 interactive performance on Llama 2 70B. In MLPerf Training v5.1, it completed a Llama 2-70B LoRA workload in ten minutes after a 2.8x generational improvement over MI300X. These results show that AMD can lead selected workloads and challenge the view that its accelerators are unsuitable for frontier deployment. They do not establish system-wide superiority, as benchmark configuration, model type, precision, networking and software tuning can all change relative performance.

The negative evidence lies in scaling and utilisation. An independent cited vLLM test reported 81% scaling efficiency for four MI300X GPUs, versus 99.8% for two H200 GPUs. An external estimate places AMD real-world model-FLOPS utilisation near 45%, versus 50–55% for Blackwell. The configurations are not comparable enough to establish a universal product ranking. However, they identify the commercial risk: accelerators must remain productive when thousands of units communicate, not just when a single node runs a curated benchmark. A 10% utilisation disadvantage can erase much of a lower purchase price because electricity, networking and facility costs continue while silicon is idle.

- Hardware evidence: MI355X delivered 119% of B200 interactive performance in the cited April 2026 Llama 2 70B inference result.
- Scaling evidence: the cited June 2026 test showed 81% efficiency for four MI300X GPUs versus 99.8% for two H200 GPUs.
- Customer evidence: MI300X qualification was reported at Microsoft Azure, Meta and Oracle, including clusters of up to 16,384 GPUs.

DATA CENTER AI ACCELERATORS FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 9,000m
MODEL ESTIMATE REPORTED EBIT	USD 1,200m
FY2025 MODELLED REVENUE	$9.0bn
2027E MODELLED EBITDA	$15.0bn
CURRENT-MARKET SOTP MULTIPLE	40.33x 2027E EBITDA
EXTERNAL 2026 SHARE ESTIMATES	5–7% to ~10%
MI355X MEMORY CAPACITY	288GB HBM3E
MI300X CLOUD RENTAL RANGE	$1.71–$6.00/GPU-hour
MODELLED 2027 ACCELERATOR REVENUE	$34.0bn
META DEPLOYMENT COMMITMENT	Up to 6GW

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.1 Data Center AI Accelerators (continued)

Business division deep dive — Profit engine.

ROCM MUST REDUCE SWITCHING COSTS FASTER THAN NVIDIA EXPANDS ITS PLATFORM

ROCm 7.14 introduced the production release of TheRock, AMD's automated open-source build and release system. The independently developed SCALE language was reported to run unmodified CUDA binaries without source-code porting. This matters because AMD competes not only with the B200 or B300 chip, but with CUDA code, libraries, developer knowledge and operating practices. An open stack can appeal to customers seeking portability and control, but openness does not ensure tested compatibility across every framework and model. The valuation assumes ROCm becomes dependable enough for high-volume inference and selected training, not that it exceeds CUDA's breadth by 2027.

Customer behaviour will determine whether software progress creates recurring economics. Qualification orders may reflect supply diversification, bargaining leverage or a memory-intensive workload. Repeat orders show whether total cost of ownership holds up in production. Character.AI and DriveNets were using MI350-series deployments by mid-2026, and one cited case study processed one billion queries per day. The next evidence needed is not another isolated benchmark. It is sustained utilisation, multi-quarter expansion and migration across more than one model family at customers that can choose NVIDIA or internal ASICs.

PRICE-PERFORMANCE OPENS THE DOOR, BUT MARGINS MUST HOLD

Public cloud prices for MI300X range from USD 1.71 to USD 6.00 per GPU-hour. MI325X quotations of USD 1.57–3.31 were estimated at roughly 40% of Blackwell's hourly price. In one Mixtral 8x7B comparison, MI300X inference cost USD 22.22 per million tokens, versus USD 28.11 for H100 SXM. These figures support an inference advantage where customers value memory capacity and cost per token over ecosystem breadth. They also suggest AMD may be gaining share partly through price, which would limit the margin convergence assumed by the market.

A cited MI300X manufacturing bill of materials is approximately USD 5,300, against estimated selling prices of USD 10,000 to USD 20,000. At the USD 15,000 midpoint, gross profit before warranty, freight, software development, customer support and allocated overhead would be USD 9,700 per unit, or 64.7%. At USD 12,500, it would be 57.6%. A USD 20,000 price would produce 73.5%. The report's 44.1% 2027 EBITDA margin allows for attractive silicon gross margin while funding software and systems investment. The reverse valuation's 66.1% group EBIT margin requires much stronger pricing, mix or operating leverage.

DATA CENTER AI ACCELERATORS FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 9,000m
MODEL ESTIMATE REPORTED EBIT	USD 1,200m
FY2025 MODELLED REVENUE	$9.0bn
2027E MODELLED EBITDA	$15.0bn
CURRENT-MARKET SOTP MULTIPLE	40.33x 2027E EBITDA
EXTERNAL 2026 SHARE ESTIMATES	5–7% to ~10%
MI355X MEMORY CAPACITY	288GB HBM3E
MI300X CLOUD RENTAL RANGE	$1.71–$6.00/GPU-hour
MODELLED 2027 ACCELERATOR REVENUE	$34.0bn
META DEPLOYMENT COMMITMENT	Up to 6GW

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.1 Data Center AI Accelerators (continued)

Business division deep dive — Profit engine.

DEMAND VISIBILITY HAS IMPROVED, BUT COMMITMENTS MUST BECOME REVENUE WITHOUT STRAINING SUPPLY

AMD and Meta announced a multi-year partnership to deploy up to 6GW of Instinct GPUs. The first 1GW, based on custom MI450 products, begins in the second half of 2026. AMD also cited 6GW of Helios commitments across customers including OpenAI, Meta and Anthropic. These announcements establish strategic demand, but should not be added together because the degree of overlap is undisclosed. Gigawatts also cannot be converted precisely into revenue without rack power, deployment timing, accelerator count and AMD content per rack.

A transparent conversion shows the range. At 1,400W accelerator TDP, 1GW divided by chip TDP alone equals about 714,000 accelerators. Real data centers also allocate power to CPUs, networking, memory, cooling and conversion losses. If accelerators represent 55–70% of rack power, the range is approximately 393,000–500,000 units per gigawatt. At USD 25,000 of AMD revenue per accelerator equivalent, one gigawatt could represent roughly USD 9.8–12.5bn of cumulative silicon revenue before custom pricing, replacement parts and deployment timing. The base case needs only part of the announced multi-year commitments to convert by 2027, but timing is highly sensitive.

PACKAGING, HBM AND EXPORT CONTROLS DETERMINE HOW MUCH DEMAND AMD CAN SERVE

TSMC targeted 120,000–130,000 monthly CoWoS wafers by end-2026, while an earlier industry estimate projected only 88,000. This report weights the later TSMC range, while recognising that AMD competes for capacity with larger customers. Reported CoWoS lead times remain 52–78 weeks, and high-end ABF substrate lead times are 52–70 weeks. AMD has reportedly committed more than USD 10bn to Taiwan packaging partnerships and secured three-year Samsung memory commitments. These actions improve supply visibility but help explain why the authoritative model forecasts negative FCFF in 2026 and 2027 despite rapid earnings growth.

Export regulation creates a separate limit. BIS moved eligible China and Macau shipments to case-by-case review below a 21,000 TPP score and a 6,500GB/s DRAM-bandwidth threshold. MI325X was specifically cited near 20,800 TPP, while MI400-class products are expected to exceed the threshold and face a presumption of denial. A further cited rule caps covered exports relative to US shipments, and a 25% tariff applies to certain non-US-produced advanced chips routed through the United States for international export. AMD recorded USD 800m of inventory charges related to earlier MI308 restrictions, although it later sold USD 390m of released MI308 inventory in Q4 2025. Case-by-case review improves access relative to blanket denial, but does not make China demand equivalent to unrestricted demand.

DATA CENTER AI ACCELERATORS FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 9,000m
MODEL ESTIMATE REPORTED EBIT	USD 1,200m
FY2025 MODELLED REVENUE	$9.0bn
2027E MODELLED EBITDA	$15.0bn
CURRENT-MARKET SOTP MULTIPLE	40.33x 2027E EBITDA
EXTERNAL 2026 SHARE ESTIMATES	5–7% to ~10%
MI355X MEMORY CAPACITY	288GB HBM3E
MI300X CLOUD RENTAL RANGE	$1.71–$6.00/GPU-hour
MODELLED 2027 ACCELERATOR REVENUE	$34.0bn
META DEPLOYMENT COMMITMENT	Up to 6GW

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.1 Data Center AI Accelerators (continued)

Business division deep dive — Profit engine.

WHAT IT IS WORTH AND WHAT BREAKS IT

The current-market SOTP assigns USD 605.0bn of EV to accelerators, equal to 40.33x the report's USD 15.0bn 2027 EBITDA estimate. The multiple exceeds NVIDIA's supplied 2027 EV/EBITDA of 15.4x because this is an allocation of AMD's actual control EV, not the report's preferred exit multiple. It shows that accelerators account for 71.7% of AMD enterprise value even after assigning meaningful value to EPYC, client and embedded. The target-price bridge tempers this market reading with a 22.0x group EBITDA method and a 30.0x earnings method.

The value holds if AMD turns qualifications into recurring multi-year deployments, reaches approximately USD 34bn of accelerator revenue, and produces about USD 15bn of EBITDA without relying on unsustainably low pricing. It fails if multi-node software performance remains materially behind, custom ASICs take more than the assumed 25% accelerator share, MI400 cannot access key geographies, or packaging and HBM constraints delay the volume ramp. At a 35% rather than 44.1% EBITDA margin, the same revenue would generate USD 11.9bn of EBITDA. Applying 30x rather than 40.33x would reduce accelerator EV to roughly USD 357bn. Conversely, USD 40bn of revenue at a 50% margin and 35x multiple would support USD 700bn of accelerator EV. Repeat deployment and margin evidence therefore matter more than an isolated market-share estimate.

DATA CENTER AI ACCELERATORS FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE
USD 9,000m

MODEL ESTIMATE REPORTED EBIT
USD 1,200m

FY2025 MODELLED REVENUE
$9.0bn

2027E MODELLED EBITDA
$15.0bn

CURRENT-MARKET SOTP MULTIPLE
40.33x 2027E EBITDA

EXTERNAL 2026 SHARE ESTIMATES
5–7% to ~10%

MI355X MEMORY CAPACITY
288GB HBM3E

MI300X CLOUD RENTAL RANGE
$1.71–$6.00/GPU-hour

MODELLED 2027 ACCELERATOR REVENUE
$34.0bn

META DEPLOYMENT COMMITMENT
Up to 6GW

AI ACCELERATOR POSITION, PRODUCT EVIDENCE AND SCALE

PLATFORM	2026 MARKET POSITION	MEMORY / PRODUCT EVIDENCE	ECONOMICS / SCALE
AMD Instinct	5–7% to ~10% est.	MI355X: 288GB HBM3E	$1.71–$6.00/GPU-hr
NVIDIA	81–85%; alt. 92%	B200: 192GB HBM3E	$193.7bn DC FY2026
Google TPU	6–8% FLOPS estimate	Internal and Cloud TPU	64.8% custom-ASIC est.
AWS Trainium	2–3% estimate	Trainium/Inferentia	14.6% custom-ASIC est.
Intel Gaudi	1–3% estimate	Gaudi 3	Minor merchant share

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.2 Data Center CPUs

Business division deep dive — Profit engine.

EPYC COMPETES FOR GENERAL-PURPOSE SERVER SPENDING AS AI CHANGES THE MIX

The global enterprise-server market is estimated at USD 101.38bn in 2026. General-purpose servers represent an estimated 30.66% of the broader data-center-server market. This is EPYC's directly contestable annual spending pool, although the market boundary is shifting. AI servers reportedly exceed 70% of server-industry value, while CPUs remain essential for orchestrating accelerators, running databases and serving general workloads. The valuation therefore assumes EPYC participates in conventional servers and accelerator hosts, rather than treating AI as pure cannibalisation. The report models USD 14.0bn of 2027 CPU revenue, equal to 13.8% of the stated 2026 enterprise-server market before market growth.

The USD 14.0bn requirement is built from AMD's reported Data Center segment, not product disclosure. FY2025 modelled CPU revenue is USD 7.0bn, with the balance of the USD 16.635bn reported Data Center segment allocated to accelerators, networking and systems integration. Doubling by 2027 requires continued unit gains, richer server mix and growth in the underlying compute pool. It does not require AMD to double its already high reported x86 revenue share, because the Q1 2026 revenue-share figure covers only x86 and excludes Arm-based CPUs.

REVENUE SHARE APPROACHES PARITY, BUT UNIT SHARE SHOWS THE REMAINING PRICE AND MIX ADVANTAGE

Mercury Research data cited by TechPowerUp places AMD at a record 46.2% of x86 server CPU revenue in Q1 2026. Another Mercury report gives AMD 33.2% of x86 server units. The 13-point gap suggests AMD's average revenue per unit exceeds the x86 market average, consistent with high-core-count server and hyperscaler strength. Q4 2025 hyperscaler revenue share was reported at 41.3%, showing that adoption extends beyond smaller enterprise buyers. This is valuable because cloud customers validate performance at scale and can deploy quickly once a platform is qualified.

Intel unit-share reports conflict. One source places Intel at 66.8% of x86 server units in Q1 2026, while another gives 54.9%; both cannot use the same denominator. This report gives greater weight to AMD's specific 33.2% x86 figure and avoids deriving Intel share by subtraction where Arm may or may not be included. The uncertainty does not change the central point: Intel remains larger in units, while AMD is much closer in revenue. The next stage of value creation depends more on retention, supply and platform cadence than on proving EPYC can compete.

- AMD: 46.2% x86 server revenue share and 33.2% x86 unit share in Q1 2026.
- Intel: reported unit-share readings range from 54.9% to 66.8% because source denominators conflict.
- Arm: 17.7% of global data-center CPU unit shipments in Q1 2026.

DATA CENTER CPUS FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 7,000m
MODEL ESTIMATE REPORTED EBIT	USD 1,000m
FY2025 MODELLED REVENUE	$7.0bn
2027E MODELLED EBITDA	$4.5bn
SOTP MULTIPLE	22.0x 2027E EBITDA
X86 SERVER REVENUE SHARE	46.2% (Q1 2026)
X86 SERVER UNIT SHARE	33.2% (Q1 2026)
ARM DATA-CENTER UNITS	17.7% (Q1 2026)
HYPERSCALER REVENUE SHARE	41.3% (Q4 2025)

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.2 Data Center CPUs (continued)

Business division deep dive — Profit engine.

ARM IS A LARGER STRUCTURAL THREAT THAN ANY SINGLE INTEL PRODUCT CYCLE

Arm-based processors reached 17.7% of data-center CPU units in Q1 2026, up from an estimated 9% of revenue by the end of 2025. AWS, which held 28% of cloud-infrastructure services in Q1 2026, has the scale to direct workloads to internal processors. Microsoft, Google and other hyperscalers are also represented in the broader Arm category, although the research register provides no separate 2026 shares for Graviton, Cobalt or Axion. This matters because internally designed CPUs can be priced on total cloud economics rather than merchant gross margin.

NVIDIA adds a second Arm threat through Grace and Vera. Its standalone and superchip CPU revenue is projected to reach USD 20bn in 2026, although this includes products sold in accelerator-centric configurations and is not directly comparable with AMD's EPYC-only model. NVIDIA can optimise CPU, GPU, networking and software as a single platform, potentially reducing the value of a standalone host CPU. AMD's response is architectural choice: EPYC can serve general workloads, pair with Instinct and attach to third-party accelerators. That flexibility must outweigh the benefits of a single-vendor stack.

EPYC ECONOMICS DEPEND ON RETAINING PREMIUM MIX WHILE INCREASING UNITS

The gap between AMD's 46.2% x86 revenue share and 33.2% x86 unit share indicates relative pricing. Dividing revenue share by unit share gives an index of 1.39 versus the market average, before adjusting for sockets, configurations and reporting definitions. This is not a disclosed ASP, but suggests AMD's current revenue is concentrated in high-value processors. If unit share rises as mix normalises, revenue may grow more slowly than shipments. Continued high-core-count leadership could preserve the premium.

The report's USD 4.5bn 2027 EBITDA estimate implies a 32.1% margin on USD 14.0bn of revenue. This is below the strongest accelerator economics and the 65% EBIT margins shown for some supplied peers. However, it reflects mature merchant-CPU competition and substantial development expense. The estimate requires competitive product cadence, foundry access to support volume and continued merchant-CPU adoption by hyperscalers. A 25% EBITDA margin would reduce EBITDA to USD 3.5bn, while a 38% margin would increase it to USD 5.3bn.

DATA CENTER CPUS FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 7,000m
MODEL ESTIMATE REPORTED EBIT	USD 1,000m
FY2025 MODELLED REVENUE	$7.0bn
2027E MODELLED EBITDA	$4.5bn
SOTP MULTIPLE	22.0x 2027E EBITDA
X86 SERVER REVENUE SHARE	46.2% (Q1 2026)
X86 SERVER UNIT SHARE	33.2% (Q1 2026)
ARM DATA-CENTER UNITS	17.7% (Q1 2026)
HYPERSCALER REVENUE SHARE	41.3% (Q4 2025)

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.2 Data Center CPUs (continued)

Business division deep dive — Profit engine.

ACCELERATOR GROWTH CAN SUPPORT EPYC, BUT ATTACHMENT IS NOT ASSURED

Every AI rack needs host processing, but an Instinct sale does not guarantee an EPYC sale. Customers can pair AMD GPUs with alternative CPUs, and AMD systems can use third-party networking. In the positive case, Helios qualification validates EPYC, Instinct and Pensando together, increasing AMD content and improving data movement. In the negative case, hyperscalers treat each component as modular and use AMD only where its price-performance is strongest.

The valuation does not assign EPYC all host-CPU revenue associated with the 6GW commitments. Instead, modelled USD 14.0bn of revenue reflects general server-share gains plus some accelerator attachment. This avoids double counting between the CPU and accelerator business divisions. Disclosure of rack bill of materials, CPU attachment rates or platform-specific revenue would reduce a major modelling uncertainty.

GEOGRAPHY ADDS OPPORTUNITY AND POLICY RISK

AMD generated USD 7.75bn of total company revenue from China in 2025, but the register does not separate EPYC exposure. Domestic Chinese CPU vendors were estimated to hold nearly 15% of their local server-CPU market in 2025, creating competition beyond Intel and Arm hyperscaler designs. CPU export constraints are generally product-specific and cannot be inferred from accelerator thresholds. The valuation therefore assumes China remains meaningful but not unrestricted, without assigning a separate regional share.

Outside China, cloud and enterprise refresh cycles remain the main drivers. North American cloud concentration supports rapid adoption, but gives a small number of buyers substantial bargaining power. If customers internalise more CPU design, AMD may retain units while merchant-system growth slows. A broader OEM and enterprise mix would improve durability, even with slower qualification cycles.

THE USD 99BN VALUE IS SUPPORTABLE WITHOUT EXTREME TERMINAL MARGINS

The SOTP assigns USD 99.0bn of EV to Data Center CPUs using USD 4.5bn of modelled 2027 EBITDA and a 22.0x multiple. The multiple is below AMD's 2026 EV/EBITDA of 59.45x and the supplied 2026 peer median of 32.3x. It is above mature semiconductor comparators such as Qualcomm and NXP. This premium reflects share gains and high-value server exposure while allowing for Arm and Intel competition. Unlike the accelerator allocation, the CPU value does not require NVIDIA-like margins.

DATA CENTER CPUS FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 7,000m
MODEL ESTIMATE REPORTED EBIT	USD 1,000m
FY2025 MODELLED REVENUE	$7.0bn
2027E MODELLED EBITDA	$4.5bn
SOTP MULTIPLE	22.0x 2027E EBITDA
X86 SERVER REVENUE SHARE	46.2% (Q1 2026)
X86 SERVER UNIT SHARE	33.2% (Q1 2026)
ARM DATA-CENTER UNITS	17.7% (Q1 2026)
HYPERSCALER REVENUE SHARE	41.3% (Q4 2025)

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.2 Data Center CPUs (continued)

Business division deep dive — Profit engine.

The base case fails if AMD's x86 unit share stops rising while revenue share falls, indicating loss of premium mix. It also fails if Arm exceeds the current 17.7% unit share without expanding the overall server pool, or if Intel improves manufacturing and product execution enough to force price competition. The upside case is a 38% EBITDA margin on USD 16bn of revenue, or USD 6.1bn of EBITDA. At 22x, this supports approximately USD 134bn of EV. The downside case is USD 11bn of revenue at a 25% margin and 16x, supporting only USD 44bn.

DATA CENTER CPUS FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE
USD 7,000m

MODEL ESTIMATE REPORTED EBIT
USD 1,000m

FY2025 MODELLED REVENUE
$7.0bn

2027E MODELLED EBITDA
$4.5bn

SOTP MULTIPLE
22.0x 2027E EBITDA

X86 SERVER REVENUE SHARE
46.2% (Q1 2026)

X86 SERVER UNIT SHARE
33.2% (Q1 2026)

ARM DATA-CENTER UNITS
17.7% (Q1 2026)

HYPERSCALER REVENUE SHARE
41.3% (Q4 2025)

SERVER CPU SHARE, ECONOMICS AND COMPETITIVE RISK

CPU PLATFORM	Q1 2026 POSITION	ECONOMIC SIGNAL	STRATEGIC RISK
AMD EPYC	46.2% x86 revenue	33.2% x86 units	Retain premium mix
Intel Xeon	54.9–66.8% units	Largest installed base	Execution recovery
Arm aggregate	17.7% DC units	9% revenue est. 2025	Internal cloud CPUs
NVIDIA Grace/Vera	$20bn 2026 forecast	GPU-platform attachment	Integrated stack
China domestic	15% local share 2025	Policy-supported supply	Regional substitution

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.3 Adaptive and Embedded Compute

Business division deep dive — Profit engine.

RYZEN COMPETES FOR A SHARE OF GLOBAL DEVICE AND CLOUD-COMPUTE SPENDING

The broadest spending pool relevant to Client Compute is the approximately USD 1.3tn global consumer-technology market projected by NIQ on 2 January 2026. PCs are only one part of that market. IDC estimated on 13 March 2026 that the global PC market would be worth USD 274bn in 2026, with higher average selling prices offsetting weaker unit volumes. The PC market is about 21.1% of the broader consumer-technology pool, calculated as USD 274bn / USD 1.3tn. Ryzen does not capture the full value of a PC, as the processor is only one component. However, AMD can influence end-system spending through CPU performance, integrated graphics, battery life and on-device AI capability. The relevant opportunity is processor and platform content within USD 274bn of annual systems, plus selected computing spending that can move from the cloud to the client device.

On-device AI creates a second displacement opportunity, not simply a new PC label. Gartner projected on 10 August 2026 that AI-optimised infrastructure-as-a-service spending could reach USD 42.2bn in 2026. It also identified potential for on-device AI compute to displace some spending as inference moves to the edge. The USD 42.2bn pool is not an AMD Client revenue forecast. Local inference requires complete devices, software and application adoption, and much cloud demand will remain centralised. It does show why buyers may value neural-processing capability if local execution reduces latency, connectivity needs, privacy risk or recurring cloud charges. Client silicon captures incremental value only when software uses the NPU and the resulting benefit affects system choice or price.

AMD HAS MATERIAL X86 SHARE, WITH REVENUE MIX MODESTLY ABOVE UNIT MIX

Mercury Research estimated on 15 May 2026 that AMD's Q1 2026 global client x86 unit share reached a record 29.6%. Mercury Research data reported by Tom's Hardware on 14 May 2026 put client x86 revenue share at 31.4% in the same quarter. The 1.8-percentage-point gap indicates that AMD's revenue mix modestly exceeded its unit mix. The report's revenue-to-unit share index is 1.061, calculated as 31.4% / 29.6%. On these measures, AMD captured about 6.1% more revenue share than unit share. The premium helps, but does not make pricing discipline optional.

ADAPTIVE AND EMBEDDED COMPUTE FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 3,500m
MODEL ESTIMATE REPORTED EBIT	USD 700m
FY2025 MODELLED REVENUE	$3.5bn
Q2 2026 REPORTED EMBEDDED REVENUE	$977m
Q2 2026 GROWTH	19% year on year
2027E MODELLED EBITDA	$1.8bn
SOTP MULTIPLE	25.0x 2027E EBITDA
AMD ADAPTIVE-MARKET AMBITION	>70% revenue share
2026 FPGA MARKET ESTIMATES	$8.9–$15.28bn

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.3 Adaptive and Embedded Compute (continued)

Business division deep dive — Profit engine.

The product evidence shows strength but also an unfinished notebook opportunity. Mercury Research estimated on 15 May 2026 that AMD held 33.2% of desktop x86 processor units in Q1 2026, down from 36.4% in Q4 2025 but higher year on year. The same source estimated AMD's mobile x86 unit share at 28.3% in Q1 2026. Desktop share was therefore 4.9 percentage points above mobile share, calculated as 33.2% minus 28.3%. Closing part of this gap through premium notebook designs would be more valuable than relying only on desktop enthusiasts, as mobile systems can offer longer platform runs and broader OEM distribution.

INTEL REMAINS THE DIRECT X86 INCUMBENT, WHILE ARM CHANGES THE MARKET BOUNDARY

Intel still controlled an estimated 70.4% of consumer x86 processor units in Q1 2026, according to Tom's Hardware using Mercury Research data on 14 May 2026. AMD's 29.6% and Intel's 70.4% shares total 100% because both refer to the x86 client market, not the total PC processor market. Arm-based processors, mainly from Apple and Qualcomm, separately reached an estimated 14.4% of total PC units in Q1 2026, according to PCWorld using Mercury Research data on 4 June 2026. The Arm figure should not be added to the x86 shares or used to infer Intel's total-PC share without reconciling the different denominators. AMD must compete both for x86 share and for x86 relevance.

- Intel held 70.4% of consumer x86 processor units in Q1 2026, according to Tom's Hardware using Mercury Research data on 14 May 2026.
- Arm-based processors, primarily from Apple and Qualcomm, represented 14.4% of total PC units in Q1 2026, according to PCWorld using Mercury Research data on 4 June 2026.
- Apple represented 9.1% of worldwide PC shipments in Q1 2026 according to Gartner on 10 April 2026, while IDC estimated 9.5% on 9 April 2026; the report uses Gartner's 9.1% figure but treats the 0.4-point difference as immaterial.
- AMD held 31.4% of client x86 processor revenue in Q1 2026, according to Mercury Research data reported by Tom's Hardware on 14 May 2026, placing it above its 29.6% unit share.

ADAPTIVE AND EMBEDDED COMPUTE FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	**USD 3,500m**
MODEL ESTIMATE REPORTED EBIT	**USD 700m**
FY2025 MODELLED REVENUE	**$3.5bn**
Q2 2026 REPORTED EMBEDDED REVENUE	**$977m**
Q2 2026 GROWTH	**19% year on year**
2027E MODELLED EBITDA	**$1.8bn**
SOTP MULTIPLE	**25.0x 2027E EBITDA**
AMD ADAPTIVE-MARKET AMBITION	**>70% revenue share**
2026 FPGA MARKET ESTIMATES	**$8.9–$15.28bn**

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.3 Adaptive and Embedded Compute (continued)

Business division deep dive — Profit engine.

Intel’s scale gives it purchasing relationships, platform familiarity and the ability to use price to defend sockets even when AMD has a performance edge. Apple is a different challenge because it controls the processor, operating system and hardware design, allowing full-platform optimisation rather than merchant CPU optimisation. Qualcomm's exact 2026 PC processor share cannot be derived from the available Arm aggregate. The 14.4% figure combines architectures and vendors, while Apple shipment estimates are system-vendor shares rather than processor-unit shares. This denominator mismatch prevents a credible subtraction for Qualcomm. The conclusion does not depend on the missing figure: Arm already has enough total-PC presence to constrain the long-term x86 opportunity.

AI-PC MONETISATION REQUIRES USEFUL SOFTWARE, NOT SIMPLY AN NPU ON THE SPECIFICATION SHEET

An NPU can increase semiconductor content and product differentiation, but only if applications perform useful work locally. The strongest use cases will likely be those where latency, privacy, power efficiency or intermittent connectivity makes cloud execution less attractive, rather than tasks moved on-device for marketing. Gartner's USD 42.2bn 2026 estimate for AI-optimised IaaS spending, published on 10 August 2026, shows a meaningful pool of remote inference expenditure to contest. Even a small shift could support higher-value client silicon, but the benefit could accrue as readily to Apple or Qualcomm as to AMD. AMD must pair competitive hardware with OEM designs and a software ecosystem that exposes NPU capability without requiring extensive application rewrites.

AI-PC adoption can also destroy value if every processor vendor offers similar capability and buyers treat the NPU as a standard feature. In that case, die area and development expense rise while system vendors resist higher processor prices. AMD's 31.4% revenue share versus 29.6% unit share in Q1 2026 offers some evidence of premium positioning. However, the report’s 1.061 share index is not large enough to prove durable AI monetisation. A successful strategy would preserve revenue share above unit share while adding notebook designs, rather than buying volume through discounts. The revenue-share premium is therefore more useful than unit share alone when assessing whether AI functionality creates economic value.

ADAPTIVE AND EMBEDDED COMPUTE FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 3,500m
MODEL ESTIMATE REPORTED EBIT	USD 700m
FY2025 MODELLED REVENUE	$3.5bn
Q2 2026 REPORTED EMBEDDED REVENUE	$977m
Q2 2026 GROWTH	19% year on year
2027E MODELLED EBITDA	$1.8bn
SOTP MULTIPLE	25.0x 2027E EBITDA
AMD ADAPTIVE-MARKET AMBITION	>70% revenue share
2026 FPGA MARKET ESTIMATES	$8.9–$15.28bn

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.3 Adaptive and Embedded Compute (continued)

Business division deep dive — Profit engine.

THE MODEL DEPENDS MORE ON PREMIUM NOTEBOOKS AND MIX DISCIPLINE THAN A BROAD UNIT RECOVERY

The report models USD 13.0bn of Client Compute revenue and USD 2.8bn of EBITDA in 2027, both as valuation estimates. The resulting EBITDA margin is 21.5%, calculated as USD 2.8bn / USD 13.0bn. Modelled revenue equals 4.7% of IDC's USD 274bn 2026 PC end-system market, calculated as USD 13.0bn / USD 274bn. This is not a processor market-share estimate because AMD revenue and complete-system spending use different value bases. The modest end-system percentage shows that the model does not require AMD to capture a large share of total PC retail value. It does require continued processor share, sufficient silicon content per system and mix that supports more than a 20% EBITDA margin.

A unit-cycle recovery alone would not validate the model if it comes from low-priced systems. Stronger premium-mobile penetration would be more constructive because AMD's Q1 2026 mobile x86 unit share of 28.3% remained below its 33.2% desktop share, according to Mercury Research on 15 May 2026. If notebook share rises without discounting, the division can increase revenue while preserving its revenue-share premium. If Intel prices aggressively or Arm systems take a disproportionate share of premium notebooks, AMD could gain units but lose mix quality. The key evidence is therefore mobile design wins, revenue share and margin, rather than headline PC shipment growth.

THE USD 44.8BN ALLOCATION IS SENSITIVE TO MARGIN AND THE DURABILITY OF X86 SHARE

The valuation applies a 16.0x multiple to the report's estimated USD 2.8bn of 2027 EBITDA, producing a USD 44.8bn EV allocation. This equals 5.31% of AMD's immutable USD 843.861bn control EV, calculated as USD 44.8bn / USD 843.861bn. The supplied peer set places Qualcomm at 10.9x 2027 EV/EBITDA, while Intel's comparison is distorted by foundry economics. The higher 16.0x multiple reflects AMD's share gains, premium mix and fabless model. It remains well below AMD's group valuation and does not value Client Compute as an AI accelerator franchise. As the division is only about one-twentieth of control EV, a strong PC cycle cannot offset a major decline in AMD's larger businesses.

The arithmetic shows the key sensitivities. Each one-turn change in the 16.0x multiple changes divisional value by USD 2.8bn, using the report's USD 2.8bn EBITDA estimate. Each 100-basis-point change in EBITDA margin on USD 13.0bn of estimated revenue changes EBITDA by USD 130m. At 16.0x, that margin movement changes indicated value by USD 2.08bn, calculated as USD 130m × 16.0. The valuation is therefore more exposed to sustained pricing and mix than to the small gap between the 9.1% and 9.5% Apple shipment estimates.

ADAPTIVE AND EMBEDDED COMPUTE FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 3,500m
MODEL ESTIMATE REPORTED EBIT	USD 700m
FY2025 MODELLED REVENUE	$3.5bn
Q2 2026 REPORTED EMBEDDED REVENUE	$977m
Q2 2026 GROWTH	19% year on year
2027E MODELLED EBITDA	$1.8bn
SOTP MULTIPLE	25.0x 2027E EBITDA
AMD ADAPTIVE-MARKET AMBITION	>70% revenue share
2026 FPGA MARKET ESTIMATES	$8.9–$15.28bn

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.3 Adaptive and Embedded Compute (continued)

Business division deep dive — Profit engine.

THE CASE FAILS IF ARM CAPTURES PREMIUM SYSTEMS OR AI FEATURES BECOME UNDIFFERENTIATED

The first risk is faster Arm adoption beyond the estimated 14.4% of total PC units reported by PCWorld and Mercury Research for Q1 2026 on 4 June 2026. This would be particularly damaging if Apple and Qualcomm-led platforms capture premium notebooks, as AMD could retain x86 share while its addressable market and revenue mix deteriorate. The second risk is an Intel response through price, platform incentives or better products, which could compress AMD's revenue-share premium even if unit share stays near 30%. The third is that AI-PC features become standardised and fail to support incremental processor pricing, leaving AMD with higher product costs but no ASP benefit. Under any of these outcomes, the report's 21.5% estimated 2027 EBITDA margin would be difficult to sustain.

The upside requires AMD to retain client x86 revenue share above 30%, keep revenue share above unit share, and narrow the 4.9-point gap between desktop and mobile unit share without sacrificing ASP. A useful confirmation would be premium notebook adoption that lifts mobile share above the Q1 2026 level of 28.3% while preserving the 31.4% overall x86 revenue share reported by Mercury Research through Tom's Hardware on 14 May 2026. Meaningful local inference could create a new reason to refresh PCs if it displaces part of Gartner's USD 42.2bn 2026 AI-optimised IaaS pool. The value accrues to AMD only if its hardware and software provide that utility more effectively than Intel, Apple or Qualcomm platforms. The investment case is therefore a share-and-mix argument in a mature market, helped but not guaranteed by the AI-PC cycle.

ADAPTIVE AND EMBEDDED COMPUTE FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 3,500m
MODEL ESTIMATE REPORTED EBIT	USD 700m
FY2025 MODELLED REVENUE	$3.5bn
Q2 2026 REPORTED EMBEDDED REVENUE	$977m
Q2 2026 GROWTH	19% year on year
2027E MODELLED EBITDA	$1.8bn
SOTP MULTIPLE	25.0x 2027E EBITDA
AMD ADAPTIVE-MARKET AMBITION	>70% revenue share
2026 FPGA MARKET ESTIMATES	$8.9–$15.28bn

ADAPTIVE AND EMBEDDED RECOVERY INDICATORS

COMPANY OR OPERATOR	RELEVANT CLIENT-COMPUTE POSITION	DATED MARKET EVIDENCE	ANALYTICAL READ-THROUGH
AMD	Ryzen x86 client processors	Q1 2026 client x86 unit share 29.6% — Mercury Research, 15 May 2026; revenue share 31.4% — Mercury Research via Tom's Hardware, 14 May 2026	Revenue share above unit share indicates a modest premium mix, with a report-calculated share index of 1.061.
Intel	Incumbent x86 client processor supplier	Q1 2026 consumer x86 unit share 70.4% — Tom's Hardware / Mercury Research, 14 May 2026	Intel retains the scale to defend sockets through pricing, platform relationships and product execution.
Apple	Vertically integrated Arm-based Mac platforms	Q1 2026 worldwide PC shipment share 9.1% — Gartner, 10 April 2026; alternative estimate 9.5% — IDC, 9 April 2026	Vertical integration makes Apple a structural premium-computing competitor rather than only a CPU vendor.
Qualcomm	Arm-based Windows PC processor competitor	Included with Apple in the 14.4% Arm-based share of total PC units in Q1 2026 — PCWorld / Mercury Research, 4 June 2026; standalone Qualcomm share was not disaggregated	The available aggregate confirms Arm pressure, but incompatible denominators prevent a credible standalone Qualcomm share calculation.

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.4 Client Compute

Business division deep dive — Profit engine.

RYZEN COMPETES WITHIN A USD 274BN PC MARKET WHERE VALUE CAN GROW DESPITE UNIT PRESSURE

IDC estimates the 2026 global PC market at USD 274bn, with higher ASPs offsetting weaker unit volumes. AMD's Q1 2026 client x86 unit share was 29.6%, while revenue share was 31.4%, indicating a modest premium mix. Desktop share was 33.2% and notebook share was 28.3%. The report models USD 13.0bn of 2027 revenue and USD 2.8bn of EBITDA. This requires only a mid-single-digit share of end-system value because processor revenue is one component of a PC.

Intel still controlled 70.4% of consumer x86 units, while Arm processors reached 14.4% of the total PC market. Apple's Q1 shipment share is estimated at 9.1–9.5%; the analysis uses Gartner's 9.1% reading and treats the difference as immaterial. Qualcomm's exact 2026 PC share is unavailable. AI PCs can raise silicon content, but may also strengthen integrated Arm platforms and do not ensure buyers pay separately for NPU capability.

THE MARGIN CASE DEPENDS ON PREMIUM NOTEBOOKS, NOT A UNIT-CYCLE RECOVERY ALONE

The USD 2.8bn EBITDA estimate implies a 21.5% margin. A cyclical shipment recovery without pricing discipline could raise revenue but add little value. A stronger premium-mobile mix could preserve AMD's revenue share above its unit share. The supplied Qualcomm 2027 EV/EBITDA is 10.9x, while Intel is distorted by foundry economics. AMD's 16.0x multiple reflects share gains and cleaner fabless economics, but remains well below the group's current multiple.

CLIENT COMPUTE FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	**USD 10,600m**
MODEL ESTIMATE REPORTED EBIT	**USD 900m**
FY2025 MODELLED REVENUE	**$10.6bn**
2027E MODELLED EBITDA	**$2.8bn**
SOTP MULTIPLE	**16.0x 2027E EBITDA**
CLIENT X86 UNIT SHARE	**29.6% (Q1 2026)**
CLIENT X86 REVENUE SHARE	**31.4% (Q1 2026)**
DESKTOP / NOTEBOOK SHARE	**33.2% / 28.3%**
ARM SHARE OF TOTAL PCS	**14.4% (Q1 2026)**

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.4 Client Compute (continued)

Business division deep dive — Profit engine.

The USD 44.8bn value fails if Arm adoption rises beyond the current 14.4% total-PC share, Intel closes performance gaps through pricing, or AI-PC features become undifferentiated. It rises if AMD sustains x86 revenue share above 30% and adds notebook designs without sacrificing ASP. Client represents only 5.3% of control EV, so even a strong PC cycle cannot offset a major accelerator de-rating.

CLIENT COMPUTE FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	**USD 10,600m**
MODEL ESTIMATE REPORTED EBIT	**USD 900m**
FY2025 MODELLED REVENUE	**$10.6bn**
2027E MODELLED EBITDA	**$2.8bn**
SOTP MULTIPLE	**16.0x 2027E EBITDA**
CLIENT X86 UNIT SHARE	**29.6% (Q1 2026)**
CLIENT X86 REVENUE SHARE	**31.4% (Q1 2026)**
DESKTOP / NOTEBOOK SHARE	**33.2% / 28.3%**
ARM SHARE OF TOTAL PCS	**14.4% (Q1 2026)**

CLIENT PROCESSOR SHARE AND COMPETITIVE POSITIONING

PLATFORM	Q1 2026 POSITION	PREMIUM SIGNAL	KEY CHALLENGE
AMD Ryzen	29.6% x86 units	31.4% x86 revenue	Notebook design wins
Intel Core	70.4% x86 units	Broad OEM coverage	Execution and pricing
Apple Silicon	9.1–9.5% PC units	Integrated premium Macs	Closed ecosystem
Qualcomm Snapdragon	Share unavailable	Power-efficient Arm	Windows compatibility

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.5 Data Center Networking and Infrastructure Processing

Business division deep dive — Profit engine.

NETWORKING COMPETES FOR THE DATA-MOVEMENT LAYER OF THE USD 653.4BN SYSTEMS BUDGET

Accelerated computing increases the value of moving data between processors, memory, storage and networks. NVIDIA reached 21.5% of data-center Ethernet switching revenue in Q1 2026 after 192.7% year-on-year growth, showing how accelerator leadership can pull networking share. SmartNICs were projected to represent 38% of the Ethernet controller and adapter market by 2026. AMD's opportunity is to attach Pensando DPUs and infrastructure processors to EPYC and Instinct deployments rather than sell isolated components.

The report models USD 2.5bn of 2027 revenue and USD 800m of EBITDA. FY2025 revenue of USD 1.2bn is a report estimate within Data Center because AMD does not disclose networking-only revenue, units or attachment rates. The 32% modelled margin assumes software and platform leverage without matching NVIDIA's integrated networking economics. The key missing KPI is the share of AMD accelerator racks and EPYC servers using AMD infrastructure processing.

CONCENTRATION CREATES BOTH A CHANNEL OPPORTUNITY AND A COMPETITIVE PROBLEM

The top five DPU vendors, NVIDIA, Intel, Marvell, AMD and Broadcom, were estimated to control 78–82% of 2025 revenue. Pensando's individual share was not found, so this report does not claim a market rank. AMD can win by bundling secure, programmable data services with compute, but customers may prefer vendors with broader switch and interconnect portfolios. The USD 17.6bn EV allocation equals 22.0x modelled 2027 EBITDA.

DATA CENTER NETWORKING AND INFRASTRUCTURE PROCESSING FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 1,200m
MODEL ESTIMATE REPORTED EBIT	USD 100m
FY2025 MODELLED REVENUE	$1.2bn
2027E MODELLED EBITDA	$800m
SOTP MULTIPLE	22.0x 2027E EBITDA
NVIDIA ETHERNET SHARE	21.5% (Q1 2026)
TOP-FIVE DPU CONCENTRATION	78–82% (2025)
SMARTNIC ADAPTER MIX	38% projected for 2026

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.5 Data Center Networking and Infrastructure Processing (continued)

Business division deep dive — Profit engine.

The value fails if attachment remains low or customers view Pensando as a niche product requiring separate software support. It rises if Helios creates a repeatable rack architecture where DPU content is standard and externally visible. Moving from 20% to 40% attachment on the same compute base could roughly double the associated addressable unit pool, although exact revenue cannot be calculated without disclosed rack content.

DATA CENTER NETWORKING AND INFRASTRUCTURE PROCESSING FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE
USD 1,200m

MODEL ESTIMATE REPORTED EBIT
USD 100m

FY2025 MODELLED REVENUE
$1.2bn

2027E MODELLED EBITDA
$800m

SOTP MULTIPLE
22.0x 2027E EBITDA

NVIDIA ETHERNET SHARE
21.5% (Q1 2026)

TOP-FIVE DPU CONCENTRATION
78–82% (2025)

SMARTNIC ADAPTER MIX
38% projected for 2026

DATA-CENTER NETWORKING AND INFRASTRUCTURE-PROCESSING POSITION

VENDOR	DECISIVE 2025–26 FIGURE	PRODUCT POSITION	AMD IMPLICATION
AMD Pensando	Share not disclosed	DPU / infrastructure	Needs rack attachment
NVIDIA	21.5% Ethernet Q1 26	Spectrum-X / BlueField	Integrated competitor
Intel	Top-five DPU vendor	IPU / adapters	CPU-channel rival
Marvell	Top-five DPU vendor	Interconnect / custom	Merchant specialist
Broadcom	$10.8bn AI semi Q2 26	Switching / custom ASIC	Scale and breadth

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.6 Radeon Discrete Graphics

Business division deep dive — Profit engine.

RADEON COMPETES IN A CONCENTRATED ADD-IN-BOARD MARKET

Global discrete desktop graphics-card shipments were 11.82m units in Q1 2026. AMD held 8%, NVIDIA 90% and Intel approximately 1%. Applying AMD's share to shipments implies roughly 0.95m quarterly Radeon units. This is the report's estimate and excludes mobile discrete GPUs and workstation products. The revised 8% Q4 2025 estimate replaced an earlier 5% reading after 400,000 desktop chips were identified.

FY2025 Radeon revenue is modelled at USD 904m after deducting semi-custom from Gaming. The division must maintain performance per dollar and software support without buying share through structurally weak gross margin. NVIDIA's scale supports a larger developer ecosystem and broader high-end portfolio, while Intel can price aggressively to build relevance.

VALUE DEPENDS ON A PROFITABLE NICHE, NOT PARITY WITH NVIDIA

The model uses USD 3.0bn of 2027 revenue and USD 500m of EBITDA, implying a 16.7% margin. This requires share recovery, stronger workstation or premium mix, and operating leverage from architectures shared with consoles and APUs. The USD 11.961bn EV allocation applies 23.922x EBITDA and represents only 1.4% of group EV.

The value fails if the 8% share cannot be sustained or software support keeps customers with NVIDIA despite higher prices. It rises if Radeon establishes a durable mid-range position and reuses Instinct software investment. As current EBIT is modelled as negative, evidence of positive standalone margin matters more than headline units.

RADEON DISCRETE GRAPHICS FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	**USD 904m**
MODEL ESTIMATE REPORTED EBIT	**USD -100m**
FY2025 MODELLED REVENUE	**$904m**
2027E MODELLED EBITDA	**$500m**
SOTP MULTIPLE	**23.922x 2027E EBITDA**
Q1 2026 AIB SHARE	**8%**
NVIDIA Q1 2026 AIB SHARE	**90%**
Q1 2026 AIB SHIPMENTS	**11.82m units**
IMPLIED AMD Q1 UNITS	**~0.95m**

DISCRETE DESKTOP GRAPHICS SHARE AND IMPLIED VOLUME

VENDOR	Q1 2026 AIB SHARE	IMPLIED QUARTERLY UNITS	COMPETITIVE POSITION
AMD Radeon	8%	~0.95m report estimate	Price-performance
NVIDIA GeForce	90%	~10.64m estimate	Scale and software
Intel Arc	~1%	~0.12m estimate	Entrant pricing
Other	~1% residual	~0.12m estimate	Niche suppliers

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.7 Semi-Custom Gaming SoCs

Business division deep dive — Profit engine.

SEMI-CUSTOM COMPETES FOR CONSOLE SILICON WITHIN A USD 25.56–30.3BN HARDWARE MARKET

Two 2026 market estimates place gaming-console hardware at USD 25.56bn and USD 30.3bn. AMD is estimated to supply nearly all high-end home-console SoCs through PlayStation 5 and Xbox Series X/S, while NVIDIA supplies the portable Switch platform. The report models FY2025 semi-custom revenue at USD 3.0bn by deducting Radeon from total Gaming, as customer revenue, units and ASPs are not disclosed.

Current-cycle economics are declining. Q2 2026 total Gaming revenue was USD 779m, down 31%, and AMD expects a significant double-digit semi-custom decline as the cycle matures. PS5 shipments fell to 1.6m in the June quarter, while Switch 2 shipped 3.82m and does not use AMD silicon. The 2027 bridge assumes USD 700m of EBITDA, balancing the trough against next-generation development payments and potential ramp preparation.

DESIGN WINS PROTECT THE NEXT CYCLE, BUT TIMING REMAINS UNCERTAIN

AMD is developing the next Xbox SoC for a 2027 launch window, and Reuters reported that AMD won the PlayStation 6 design over Intel. Memory inflation may push launches toward 2027–28, while an estimated PS6 bill of materials approaches USD 900. Delays defer AMD revenue, and high system costs can reduce launch volumes even when the silicon design win is secure.

The USD 10.5bn value applies 15.0x modelled 2027 EBITDA. It fails if either launch slips materially, customer volumes disappoint, or custom pricing does not cover advanced-node and memory-interface development. It rises if AMD retains both major home-console platforms and expands semi-custom designs into handheld or hybrid systems.

SEMI-CUSTOM GAMING SOCS FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 3,000m
MODEL ESTIMATE REPORTED EBIT	USD 194m
FY2025 MODELLED REVENUE	$3.0bn
Q2 2026 TOTAL GAMING REVENUE	$779m
Q2 GAMING DECLINE	31% year on year
2027E MODELLED EBITDA	$700m
SOTP MULTIPLE	15.0x 2027E EBITDA
NEXT XBOX TARGET	2027 launch window
PLAYSTATION 6 SUPPLIER	AMD design selected

CONSOLE AND HANDHELD SEMI-CUSTOM POSITIONING

PLATFORM	LATEST DATED VOLUME	SILICON SUPPLIER	NEXT-CYCLE POSITION
PlayStation 5	1.6m Q2 2026	AMD	PS6 design awarded
Xbox Series X/S	Units undisclosed	AMD	Project Helix 2027
Nintendo Switch 2	3.82m Q2 2026	NVIDIA	AMD not supplier
Valve Steam Deck	48% tracked niche 2024	AMD	Handheld opportunity
Intel handheld	Benchmark claim 2026	Intel	New integrated rival

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.8 AI Infrastructure Systems and Integration

Business division deep dive — Profit engine.

THE SYSTEMS OPPORTUNITY IS THE RACK AND DEPLOYMENT SPENDING AROUND EACH ACCELERATOR

The broad 2026 data-center-systems pool is USD 653.4bn, while IDC's differently defined 2025 AI-infrastructure spending estimate is USD 334bn. A narrower market-study definition gives only USD 58.78bn for 2025, showing that market labels differ according to whether they include facilities, networking, storage and services. This business division competes for rack design, integration and enablement spending attached to AMD silicon. It does not capture the full infrastructure bill. The report models USD 3.0bn of 2027 revenue and USD 500m of EBITDA, requiring less than 1% of the broad spending pool.

AMD's strategic role is to reduce deployment friction, not become a low-margin contract manufacturer. The sale of ZT Systems' manufacturing business for USD 3bn in cash and stock indicates a deliberate separation between design and enablement capabilities and balance-sheet-intensive production. Revenue booking between Data Center and acquisition-related lines is not separately disclosed. The USD 435m FY2025 allocation is therefore the report's estimate, not reported revenue. It recognises that integration work earns revenue while avoiding the assumption that all accelerator-related rack revenue belongs to AMD.

QUALIFICATION AT SCALE IS THE DECISIVE GATE

The 6GW of announced Helios commitments signals demand, but qualification must cover thermals, power delivery, networking, firmware, serviceability and software. A 1,400W accelerator cannot be assessed separately from liquid cooling and rack-level power density. Supermicro is estimated to hold 70–80% of the direct-liquid-cooled AI-rack niche, while Dell held 46.7% of the broader 2025 OEM AI-server market. AMD must complement these operators rather than disrupt the channels that deploy its silicon.

AI INFRASTRUCTURE SYSTEMS AND INTEGRATION FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE	USD 435m
MODEL ESTIMATE REPORTED EBIT	USD -300m
FY2025 MODELLED REVENUE	$435m
2027E MODELLED EBITDA	$500m
SOTP MULTIPLE	20.0x 2027E EBITDA
HELIOS COMMITMENTS	6GW during 2026–27
AI-SERVER LEADER	Dell 46.7% (2025)
DLC RACK LEADER ESTIMATE	Supermicro 70–80% (2026)

CORE ANALYSIS · BUSINESS DIVISION DEEP DIVES

08.8 AI Infrastructure Systems and Integration (continued)

Business division deep dive — Profit engine.

The value fails if customers qualify Instinct cards but standardise on third-party or internally designed rack architectures that leave AMD with little integration revenue. It also fails if system responsibility creates warranty and support obligations without adequate pricing. The USD 10.0bn EV allocation uses 20.0x modelled 2027 EBITDA, reflecting strategic scarcity but far less value than accelerator silicon. Multiple production Helios deployments, disclosed systems revenue and positive service margins would support a higher allocation.

AI INFRASTRUCTURE SYSTEMS AND INTEGRATION FACT SHEET

PROFIT ENGINE

MODEL ESTIMATE REVENUE
USD 435m

MODEL ESTIMATE REPORTED EBIT
USD -300m

FY2025 MODELLED REVENUE
$435m

2027E MODELLED EBITDA
$500m

SOTP MULTIPLE
20.0x 2027E EBITDA

HELIOS COMMITMENTS
6GW during 2026–27

AI-SERVER LEADER
Dell 46.7% (2025)

DLC RACK LEADER ESTIMATE
Supermicro 70–80% (2026)

AI SYSTEMS QUALIFICATION AND CHANNEL CONTEXT

OPERATOR	MARKET POSITION	DEPLOYMENT CAPABILITY	AMD RELATIONSHIP
AMD / Helios	Share not disclosed	6GW commitments	Platform designer
Dell	46.7% AI OEM (2025)	Global enterprise channel	Potential OEM partner
Supermicro	25.0% AI OEM (2025)	70–80% DLC est. 2026	Potential OEM partner
HPE	11.7% AI OEM (2025)	Enterprise integration	Potential OEM partner
Lenovo	1.7% AI OEM (2025)	Global server channel	Potential OEM partner

CORE ANALYSIS · CROSS-DIVISION BRIDGE

09 Putting the divisions back together

How the business division valuations aggregate to the group, and the re-rating the target price requires.

EPYC, Instinct, Pensando and systems integration share architecture, software and customer qualification. A successful Helios rack can increase content per deployment. The analysis values each component separately to avoid assigning the same rack revenue to multiple business divisions.

Client and Radeon share CPU, GPU and software development, while semi-custom volumes help spread architecture costs across console generations. This improves group economics. However, aggressive Radeon pricing or console customisation can consume engineering resources without generating accelerator-like margins.

Adaptive and Embedded Compute diversifies customer duration and can support long-cycle platform investment. The cash bridge is currently strained: authoritative FCFF is negative USD 10.5bn in 2026 and USD 6.3bn in 2027. Supply commitments therefore turn strategic opportunity into financing and execution risk.

BRIDGE MECHANICS

The current-market SOTP allocates the immutable USD 843.861bn EV across eight business divisions. The target then blends this market allocation with group EBITDA and earnings methods that require less extreme terminal economics.

AMD trades at 59.45x 2026 EV/EBITDA and 126.36x earnings. The supplied 2026 peer medians are 32.3x and 38.0x, respectively.

The reverse valuation requires USD 52.5bn of 2027 EBIT, versus USD 26.5bn in the followed model. Its 66.1% EBIT margin is closer to leading platform peers than to AMD’s demonstrated group margin.

NVIDIA’s supplied 2027 EV/EBITDA is 15.4x and Broadcom’s is 27.7x. AMD’s premium may be justified by faster forecast growth, but only if the investment phase produces high-margin cash rather than ongoing capacity absorption.

The target retains a 40% weight for the market SOTP, recognising real accelerator upside. The lower group methods reject the assumption that every qualification and commitment converts into NVIDIA-like economics on schedule.

CORE ANALYSIS · SCENARIOS & VERDICT

10 Scenario logic and the bottom-line judgment

How the conclusion changes under less favourable scenarios.

Bear: 2027 revenue reaches USD 65.0bn and EBITDA reaches USD 15.0bn, a 23.1% margin. Accelerator deployments slip, pricing stays aggressive, EPYC share stabilises rather than expands, and working-capital commitments constrain cash conversion. An 18.0x implied group multiple reflects attractive growth but less platform scarcity.

Base: revenue follows the authoritative USD 79.383bn forecast and EBITDA reaches USD 26.613bn, a 33.5% margin. MI400 enters scaled deployment, ROCm is production-capable for selected workloads, EPYC remains strong and Embedded normalises. The 31.72x implied SOTP multiple is the current-market allocation, not the sole target method.

Bull: revenue reaches USD 90.0bn and EBITDA reaches USD 36.0bn, a 40.0% margin. Multi-gigawatt commitments convert quickly, accelerator pricing holds, systems and networking attach at high rates, and supply spending generates positive forward cash flow. A 38.0x group multiple values AMD as a durable second AI platform.

Net debt should vary by outcome. The bear case would likely require more financing than the base USD 15.1bn 2027 model balance. The bull case should deliver better cash conversion and lower net debt after the investment peak. The scenario engine should therefore not treat identical equity bridges as an operating assumption.

BOTTOM LINE

AMD has credible products and customer access across every critical compute layer. However, 71.7% of current EV is attributed to accelerators, and the reverse valuation requires a 66.1% 2027 EBIT margin. The evidence supports a valuable second platform, but not yet the margin certainty priced into USD 514.39. A USD 403.84 target preserves substantial AI value while compensating investors for software, supply, export and cash-conversion risk.

FINANCIAL BRIDGE · PART 1

Net sales, EBITDA / EBIT, EBIT margin

14

Group-level trajectory for sales, earnings and margin; shaded bands mark forecast periods.

Net sales

USD millions

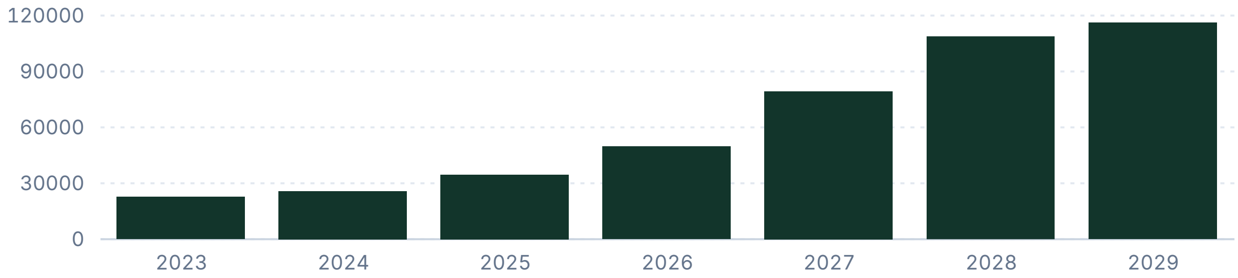

EBITDA & EBIT

USD millions

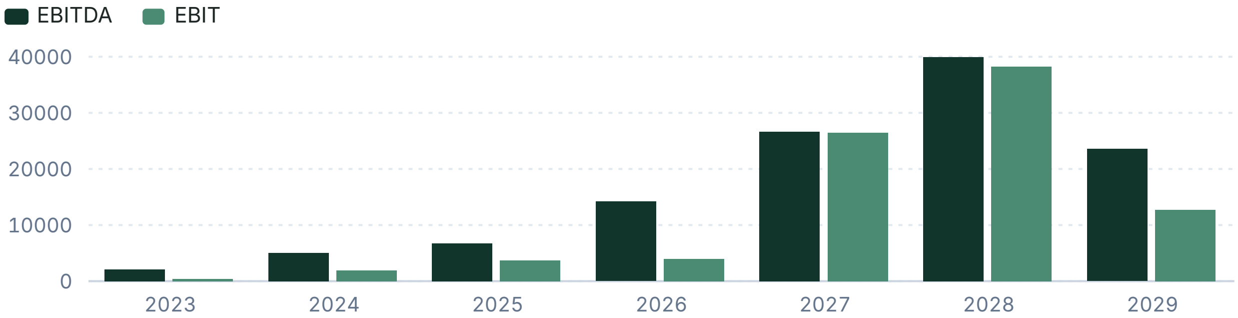

EBIT and EBITDA margin

% of net sales

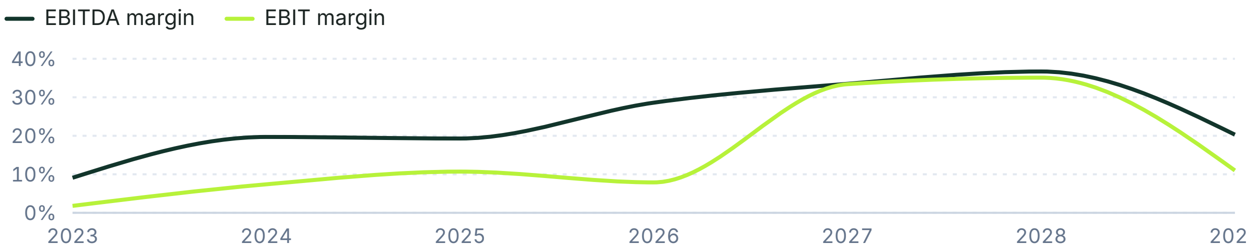

FINANCIAL BRIDGE · PART 2

15

Free cash flow, leverage, and key financials

Cash generation and balance-sheet capacity.

Free cash flow

USD millions

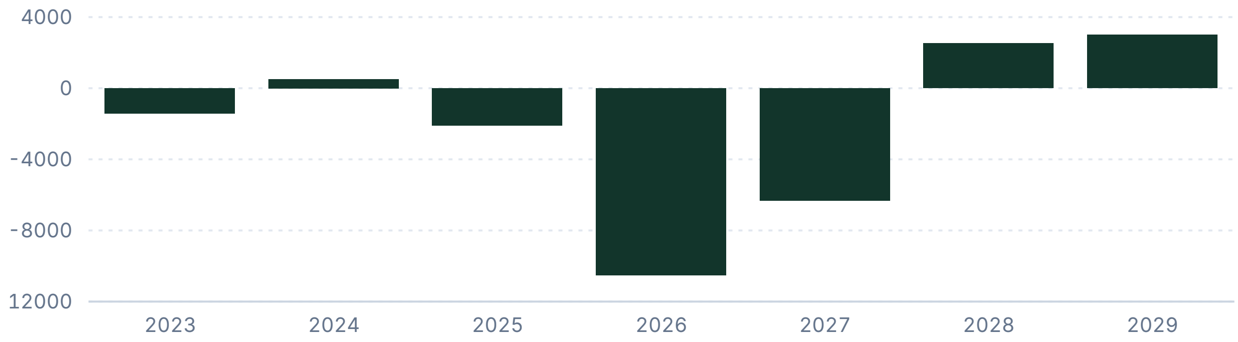

Net debt / EBITDA

Ratio

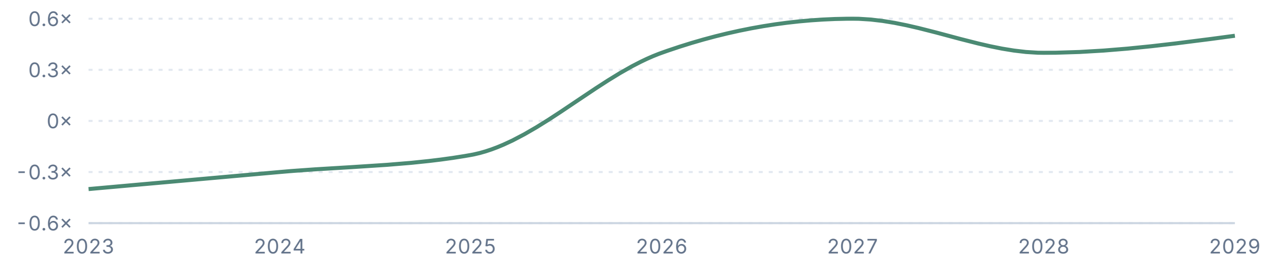

KEY FINANCIALS

	2025	2026	2027	2028	2029
Net Sales	34,639	49,714	79,383	108,870	116,229
EBITDA	6,698	14,194	26,613	39,971	23,565
Comparable EBIT	3,694	3,934	26,492	38,225	12,740
Net Earnings	**4,335**	**6,611**	**15,918**	**25,381**	**9,325**
EPS (USD)	2.7	4.1	9.8	15.6	5.7
DPS (USD)	0	0	0	0	0

VALUATION

Target price derivation

16

The rationale for the chosen valuation method and multiple, with the company's historical multiples.

SOTP is necessary because accelerator, CPU, embedded, client and gaming economics differ materially. The current-market SOTP uses division EBITDA and multiples to allocate the immutable control EV; it receives 40% weight because it transparently preserves the market's AI option value. A 22.0x 2027E EV/EBITDA method receives 35% and 30.0x 2027E P/E receives 25%, both within supplied peer and historical brackets. Sensitivity shows that EBITDA delivery and the retained multiple are equally important.

What the market is pricing.

Non-accelerator pools receive USD 238.861bn of control EV, leaving USD 605.0bn for accelerators. Against the report's USD 15.0bn 2027 accelerator EBITDA, the current price therefore implies 40.33x. At a 30x platform multiple, the market instead requires USD 20.2bn of accelerator EBITDA—USD 5.2bn above the base estimate—or equivalent systems economics.

The market assumes AMD can approach the reverse valuation's 66.1% 2027 EBIT margin. We assign a different reading: The base followed model forecasts 33.4%; retaining only part of the implied margin premium contributes directly to the lower target. (Valuatum reverse valuation and authoritative forecast model, control period 2026)

The market assumes Accelerator qualifications will convert into broad, high-utilisation production without a lasting software discount. We assign a different reading: Selected benchmarks are competitive, but cited multi-GPU scaling of 81% versus 99.8% for H200 supports a continuing execution discount. (AIMultiple, 30 June 2026; MLCommons via AMD Engineering, 1 April 2026)

The market assumes High-end accelerator demand is globally addressable on similar economics. We assign a different reading: MI325X requires case-by-case China review and MI400-class products are expected above denial thresholds, limiting geographic fungibility and raising inventory risk. (BIS framework via ECCN Finder and Tom's Hardware, January–April 2026)

HISTORICAL AND FORECAST MULTIPLES

METRIC	2020A	2021A	2022A	2023A	2024A	2025A	2026E	NORM. AVG. 2020-2025 (OUTLIERS DROPPED)
P/E	43.66x	56.06x	76.60x	278.59x	123.59x	80.54x	126.36x	76.09x
P/E adj.	44x	56x	77x	279x	124x	81x	126x	76x
EV/EBITDA	62.47x	42.66x	17.94x	114.32x	39.64x	51.97x	59.45x	42.94x
EV/EBIT	78.63x	48.08x	78.43x	591.00x	105.91x	94.23x	214.51x	81.05x
EV/Sales	11.03x	10.67x	4.20x	10.45x	7.80x	10.05x	16.97x	9.03x
P/FCF	-848.31x	77.45x	-2.21x	-167.65x	396.89x	-167.24x	-79.49x	237.17x
P/BV	18.62x	23.65x	1.85x	4.26x	3.52x	5.54x	11.99x	3.79x
P/S	11.13x	10.79x	4.28x	10.49x	7.87x	10.08x	16.80x	9.11x

VALUATION

Peers and target price bridge

17

Validated peer multiples, each method's implied price and weight, and the sensitivity of the target.

PEER COMPARISON

COMPANY	2026E EV/EBITDA	2027E EV/EBITDA	2026E EBIT MARGIN
AMD	59.45x	31.71x current EV	7.9%
NVIDIA	34.2x	15.4x	60.4%
Broadcom	32.3x	27.7x	65.7%
Intel	45.6x	49.8x	12.3%
Marvell	26.5x	102.6x	16.3%
Qualcomm	10.3x	10.9x	33.1%
Arm	94.0x	79.0x	42.5%
Peer median	32.3x	21.6x	23.4%

TARGET PRICE BRIDGE

METHOD	FORECAST METRIC	METRIC VALUE	SELECTED MULTIPLE	IMPLIED PRICE (12M DISC.)	WEIGHT
· Data Center AI Accelerators	2027E Modelled EBITDA	15,000	40.3x	371.0 USD	
· AI Infrastructure Systems and Integration	2027E Modelled EBITDA	500	20.0x	6.1 USD	
· Data Center CPUs	2027E Modelled EBITDA	4,500	22.0x	60.7 USD	
· Data Center Networking and Infrastructure Processing	2027E Modelled EBITDA	800	22.0x	10.8 USD	
· Client Compute	2027E Modelled EBITDA	2,800	16.0x	27.5 USD	
· Adaptive and Embedded Compute	2027E Modelled EBITDA	1,800	25.0x	27.6 USD	
· Semi-Custom Gaming SoCs	2027E Modelled EBITDA	700	15.0x	6.4 USD	
· Radeon Discrete Graphics	2027E Modelled EBITDA	500	23.9x	7.3 USD	
Sum of the parts	2027E Modelled division ...	n/a	n/a	514.4 USD	40%
EV/EBITDA	2027E EBITDA	26,613	22.0x	355.9 USD	35%
P/E	2027E EPS	9.8	30.0x	294.0 USD	25%
Weighted target price				403.8 USD	100%

SENSITIVITY — IMPLIED SHARE PRICE

2027E EBITDA	18.0X	20.0X	22.0X	24.0X	26.0X
21,290.4 (-20%)	231.9	258.0	284.1	310.2	336.4
23,951.7 (-10%)	261.3	290.7	320.0	349.4	378.8

VALUATION MAP

ROE vs P/BV — pricing versus profitability

18

The whole Wisdom universe as context: how the market prices profitability on average, and what the company's forecast position would imply for the share price.

ROE 2025 vs P/BV 2025

Advanced Micro Devices, Inc. in the Wisdom universe (4 632 companies)

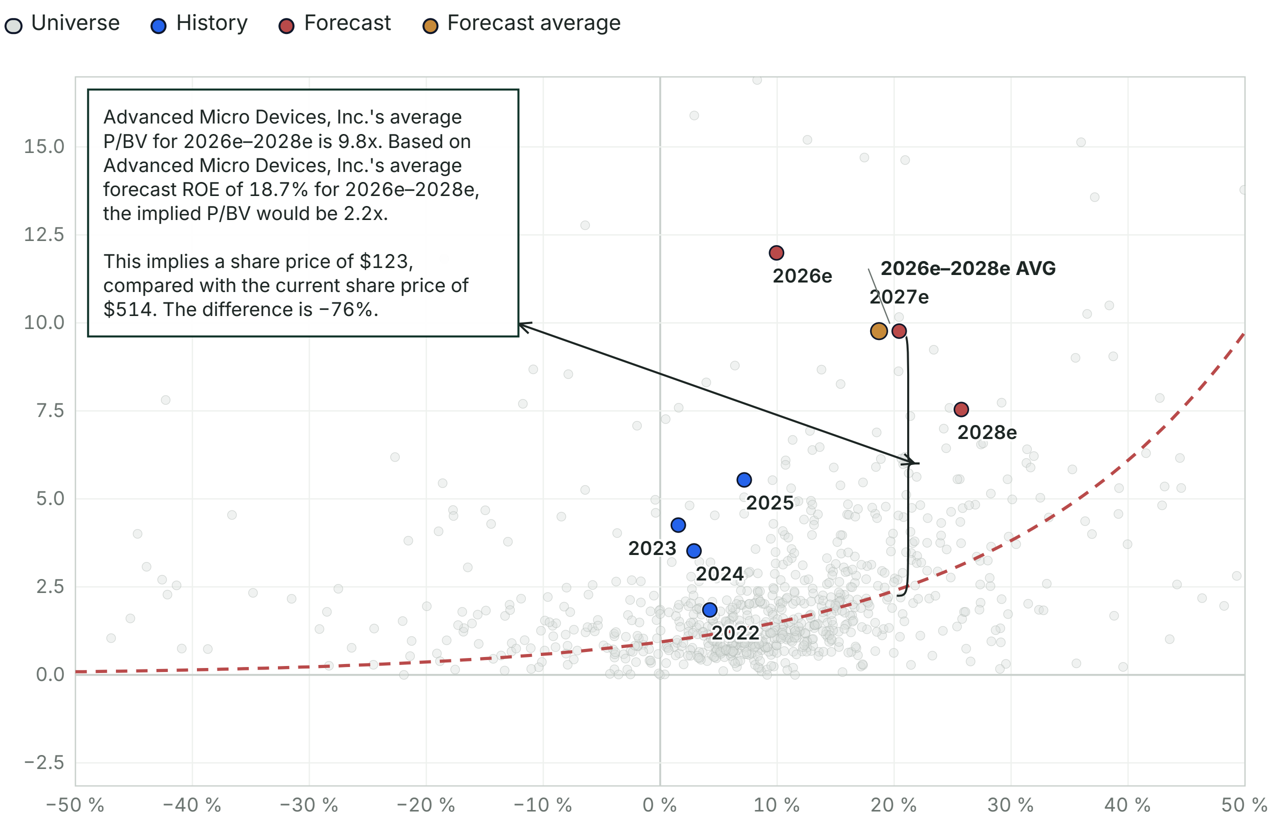

AMD's average 2026–28 forecast ROE is 18.7%, while its average P/BV is 9.8x. The universe curve implies 2.2x P/BV at that return. Applying 2.2x to average forecast book value per share of USD 54.6 gives USD 122.8, 76% below the current price. The scatter is not a target because it excludes intangible platform value and growth, but it confirms that balance-sheet returns alone cannot support the valuation.

MULTIPLES & SENSITIVITY

19

Forward multiples and EBITDA × multiple grid

How the implied valuation responds to changes in EBITDA and the applied multiple.

FORWARD MULTIPLES	2026
P/E	126.4×
EV/EBITDA	59.5×
EV/EBIT	214.5×
P / Valuatum FOCF	-79.5×
P/BV	12.0×
Dividend Yield	—
Net Debt / EBITDA	0.4×

EBITDA × EV/EBITDA sensitivity

2027E EBITDA (USDM) EV/EBITDA →	18.0×	20.0×	22.0×	24.0×	26.0×
−20% **USD 11,355m**	USD 204,390m USD 122.20 / sh	USD 227,100m USD 136.10 / sh	USD 249,810m USD 150.10 / sh	USD 272,520m USD 164.00 / sh	USD 295,230m USD 177.90 / sh
−10% **USD 12,775m**	USD 229,950m USD 137.90 / sh	USD 255,500m USD 153.60 / sh	USD 281,050m USD 169.20 / sh	USD 306,600m USD 184.90 / sh	USD 332,150m USD 200.60 / sh
Base **USD 14,194m**	USD 255,492m USD 153.60 / sh	USD 283,880m USD 171.00 / sh	USD 312,268m USD 188.40 / sh	USD 340,656m USD 205.80 / sh	USD 369,044m USD 223.20 / sh
+10% **USD 15,613m**	USD 281,034m USD 169.20 / sh	USD 312,260m USD 188.40 / sh	USD 343,486m USD 207.50 / sh	USD 374,712m USD 226.70 / sh	USD 405,938m USD 245.80 / sh
+20% **USD 17,033m**	USD 306,594m USD 184.90 / sh	USD 340,660m USD 205.80 / sh	USD 374,726m USD 226.70 / sh	USD 408,792m USD 247.60 / sh	USD 442,858m USD 268.50 / sh

20

SCENARIO VALUATION

Bear / Base / Bull bridge

Revenue, EBITDA, multiple, EV, equity, value-per-share and upside in each scenario.

SCENARIO BRIDGE

	REVENUE	EBITDA	MARGIN	MULTIPLE	EV	EQUITY	VALUE / SHARE	UPSIDE
Bear	65,000	15,000	23.1%	18.0×	270,000	264,904	USD 162.46	−68.4%
Base	79,383	26,613	33.5%	31.7×	844,164	839,068	USD 514.58	0.0%
Bull	90,000	36,000	40.0%	38.0×	1,368,000	1,362,904	USD 835.83	+62.5%

KEY CONDITIONS

The bear case reflects delayed MI400 conversion, lower accelerator utilisation and continued negative cash absorption. Base requires the authoritative revenue path, USD 15bn of accelerator EBITDA and stable EPYC premium mix. Bull requires rapid conversion of multi-gigawatt commitments, higher rack attachment, at least 40% group EBITDA margin and materially better cash conversion. Success therefore cannot use unchanged base-case net debt. These are operating boundary cases, not probability-weighted target-price inputs. The code-verified 12-month target of USD 403.8 lies between Bear and Base: Bear USD 162.5, Base USD 514.6, Bull USD 835.8. Its method weights are shown in the Target Price Bridge.

THESIS BREAKER

Sustained group EBITDA margins above 40%, repeat accelerator orders and positive post-investment FCFF would invalidate the downside framework.

21

CATALYSTS

Monitoring checklist

Near-, medium- and long-term events that could change the recommendation.

CATALYST REGISTER

	TIMING	SEGMENT	INDICATOR	VALUATION IMPACT
First material MI450 revenue from Meta deployment	NEAR-TERM	Data Center AI Accelerators	Revenue, installed MW and repeat orders	Validates conversion of the 6GW commitment
Production Helios rack qualification at multiple customers	NEAR-TERM	AI Infrastructure Systems and Integration	Named deployments and rack acceptance	Supports systems revenue and attachment
ROCm multi-node performance closes the utilisation gap	MEDIUM-TERM	Data Center AI Accelerators	Independent scaling and MFU benchmarks	Raises sustainable margin and share
EPYC sustains over 45% x86 server revenue share	NEAR-TERM	Data Center CPUs	Mercury revenue and unit share	Protects the USD 99bn CPU allocation
Embedded recovery extends beyond channel restocking	MEDIUM-TERM	Adaptive and Embedded Compute	Orders, distributor inventory and growth	Supports premium embedded multiple
Project Helix and PlayStation 6 production schedules firm	LONG-TERM	Semi-Custom Gaming SoCs	Tape-out, launch timing and customer guidance	Bridges the console-cycle trough
Pensando becomes standard content in Helios deployments	MEDIUM-TERM	Data Center Networking and Infrastructure Processing	Disclosed attachment rate or platform count	Makes networking revenue scalable

READ-THROUGH

Near-term upside could come from MI450 deployment and continued EPYC share gains. Embedded and next-generation console milestones offer smaller diversification benefits. A catalyst that raises revenue but worsens working capital or requires price concessions has less valuation impact than the headline implies.

RISKS

Risk register

22

Each risk's business division, mechanism, early-warning trigger, impact band and structural type.

RISK REGISTER

	SEGMENT	MECHANISM	EARLY WARNING	IMPACT	TYPE
ROCm remains costly to deploy at multi-node scale	AI Infrastructure Systems and Integration	Lower utilisation offsets hardware price advantages	Scaling efficiency remains materially below CUDA	HIGH	STRUCTURAL
MI400 export access remains restricted	Data Center AI Accelerators	High-end products cannot address material China demand	Presumption-of-denial decisions or inventory charges	HIGH	MANAGEABLE
Advanced packaging or HBM supply slips	Data Center CPUs	Revenue moves right while commitments consume cash	Lead times stay above one year or inventories rise	HIGH	THESIS BREAKER
Custom ASICs take more accelerator spending	Data Center Networking and Infrastructure Processing	Hyperscalers internalise silicon and system economics	TPU, Trainium or MTIA deployments accelerate	HIGH	STRUCTURAL
Arm erodes merchant server CPU demand	Client Compute	Cloud and PC vendors redirect workloads to internal Arm	Arm share rises materially above 17.7% in servers	HIGH	STRUCTURAL
Capacity investment fails to convert to FCFF	Adaptive and Embedded Compute	Working capital and supply prepayments absorb earnings	FCFF stays negative after the 2027 ramp	HIGH	THESIS BREAKER
Next-generation console launches slip	Semi-Custom Gaming SoCs	Cycle trough lasts longer and development payback defers	2027 schedules move into 2028 or later	MEDIUM	MANAGEABLE
Radeon share gains require uneconomic pricing	Client Compute	Volume grows without adequate gross profit	AIB share rises while Gaming profit deteriorates	MEDIUM	MANAGEABLE

READ-THROUGH

Diversified business divisions reduce operating volatility, but not valuation concentration. Client, Embedded, gaming and Radeon together cannot replace hundreds of billions of accelerator EV. The earliest portfolio-level warning would be rapid Data Center revenue growth alongside weak EBITDA or FCFF conversion.

23

APPENDIX

Income statement

INCOME STATEMENT — USD MILLIONS	2023	2024	2025	2026 *E*	2027 *E*	2028 *E*
Net Sales	**22,680**	**25,785**	**34,639**	*49,714*	*79,383*	*108,870*
EBITDA	**2,073**	**5,077**	**6,698**	*14,194*	*26,613*	*39,971*
EBITDA margin	9.1%	19.7%	19.3%	*28.6%*	*33.5%*	*36.7%*
Depreciation	-1,672	-3,177	-3,004	*-10,260*	*-121*	*-1,746*
Operating Profit (EBIT)	**401**	**1,900**	**3,694**	*3,934*	*26,492*	*38,225*
EBIT margin	1.8%	7.4%	10.7%	*7.9%*	*33.4%*	*35.1%*
Net financial items	107	122	472	*6,048*	*-4,981*	*-3,927*
Pre-tax Profit	**508**	**2,022**	**4,166**	*9,982*	*21,511*	*34,298*
Net Earnings	**854**	**1,641**	**4,335**	*6,611*	*15,918*	*25,381*
PER SHARE						
EPS	0.5	1	2.7	*4.1*	*9.8*	*15.6*
DPS	0	0	0	*0*	*0*	*0*
Payout ratio	—	—	—	—	—	—

E = Valuatum estimate (not yet actualised)

24

APPENDIX

Balance sheet

BALANCE SHEET — USD MILLIONS	2023	2024	2025	2026 *E*	2027 *E*	2028 *E*
ASSETS						
Tangible assets	2,222	2,425	2,312	*3,793*	*6,057*	*8,307*
Intangibles	21,363	18,930	16,705	*23,975*	*38,283*	*52,504*
Goodwill	24,262	24,839	25,126	*25,126*	*25,126*	*25,126*
Non-current assets	**50,751**	**49,489**	**49,595**	***58,346***	***74,918***	***91,389***
Inventories	4,351	5,734	7,920	*11,258*	*17,976*	*24,654*
Receivables	8,484	9,528	13,488	*17,743*	*24,050*	*30,320*
Cash & equivalents	3,933	3,787	5,539	*7,723*	*12,332*	*16,912*
Current assets	**16,768**	**19,049**	**26,947**	***36,723***	***54,358***	***71,886***
Total Assets	**67,885**	**69,226**	**76,926**	***95,453***	***129,661***	***163,658***
EQUITY & LIABILITIES						
Equity	55,892	57,568	62,999	*69,951*	*85,869*	*111,249*
Long-term debt	1,717	1,721	2,973	*6,023*	*13,249*	*15,673*
Short-term debt	751	—	874	*6,023*	*13,249*	*15,673*
Long-term liabilities	**3,567**	**3,537**	**3,534**	***6,584***	***13,810***	***16,234***
Current liabilities	**6,689**	**7,281**	**9,455**	***17,833***	***28,734***	***34,810***
Total liabilities & equity	**67,885**	**69,226**	**76,926**	***95,453***	***129,661***	***163,658***
CAPITAL STRUCTURE & RATIOS						
Net debt	-930	-1,575	-1,067	*5,096*	*15,102*	*15,485*
Capital invested	54,427	55,502	61,307	*74,275*	*100,035*	*125,683*
Equity ratio	82.3%	83.2%	81.9%	*73.3%*	*66.2%*	*68.0%*
Gearing	-1.7%	-2.7%	-1.7%	*7.3%*	*17.6%*	*13.9%*
Net debt / EBITDA	-0.4×	-0.3×	-0.2×	*0.4×*	*0.6×*	*0.4×*
Current ratio	2.5	2.6	2.9	*2.1*	*1.9*	*2.1*

E = Valuatum estimate (not yet actualised)

25

APPENDIX

Cash flow

CASH FLOW — USD MILLIONS	2023	2024	2025	2026 *E*	2027 *E*	2028 *E*
OPERATIONS						
Cash from operations (model)	-566	3,734	2,493	*12,508*	*6,688*	*17,832*
Operating cash flow (Valuatum calculation)	-1,685	2,460	2,211	*8,502*	*10,374*	*20,738*
Change in working capital	3,082	1,084	4,846	*4,364*	*9,352*	*9,294*
INVESTING & FREE CASH						
Gross capex	-80	1,915	3,110	*19,011*	*16,693*	*18,216*
Capex (ex. M&A)	80	-1,915	-3,110	*-19,011*	*-16,693*	*-18,216*
Cash after capex (CFO - gross capex)	**-646**	**1,819**	**-617**	***-6,504***	***-10,005***	***-384***
Free operating cash flow (Valuatum def.)	**-1,419**	**511**	**-2,154**	***-10,509***	***-6,319***	***2,522***
Free cash flow to firm	**-1,419**	**511**	**-2,088**	***-10,509***	***-6,319***	***2,522***
FINANCING						
CF from financing	614	-790	2,101	*8,687*	*14,614*	*4,965*
Dividends paid	—	—	—	—	—	—
Net change in cash	**128**	**1,029**	**1,484**	***2,184***	***4,609***	***4,581***

E = Valuatum estimate (not yet actualised)

26

APPENDIX

Quarterly snapshot

HISTORICAL QUARTERS — LAST TWO COMPLETED YEARS, USD MILLIONS

	Q1/24	Q2/24	Q3/24	Q4/24	Q1/25	Q2/25	Q3/25	Q4/25
Net Sales	5,473	5,835	6,819	7,658	7,438	7,685	9,246	10,270
EBITDA	820	1,038	1,562	1,657	1,548	623	2,024	2,503
EBITDA margin	15.0%	17.8%	22.9%	21.6%	20.8%	8.1%	21.9%	24.4%
Comparable EBIT	36	269	724	871	806	-134	1,270	1,752
EBIT margin	0.7%	4.6%	10.6%	11.4%	10.8%	-1.7%	13.7%	17.1%
Net Earnings	123	265	771	482	709	872	1,243	1,511
EPS	0.1	0.2	0.5	0.3	0.4	0.5	0.8	0.9

CURRENT YEAR — QUARTERLY, USD MILLIONS

	Q1/26	Q2/26	Q3/26 *E*	Q4/26 *E*
Net Sales	10,253	11,536	*12,400*	*15,525*
EBITDA	2,233	2,755	*5,492*	*7,598*
EBITDA margin	21.8%	23.9%	*44.3%*	*48.9%*
Comparable EBIT	1,476	1,990	*1,703*	*2,648*
EBIT margin	14.4%	17.3%	*13.7%*	*17.1%*
Net Earnings	1,383	2,297	*2,306*	*4,512*
EPS	0.9	1.4	*1.4*	*2.8*

E = Valuatum estimate (not yet actualised)

27

SOURCES & DISCLAIMER

Sources, assumptions, and standard disclaimer

SOURCES REGISTER

	VALUE	CATEGORY	USED IN
Financial statements and forecast model	Valuatum Wisdom model 5639, snapshot period 2026	FINANCIAL SNAPSHOT	Financials, forecasts & valuation
Share-price history and market profile	Financial Modeling Prep, AMD financialmodelingprep.com	MARKET DATA	Price chart & market facts
Form 10-K/A (Amendment No. 1) - Annual Report for the fisca...	ir.amd.com/.../0000002488-26-000021.pdf	OFFICIAL SOURCE	Reported segment table
NVIDIA's FY2026 data center segment revenue reached a recor...	vertexaisearch.cloud.google.com	OFFICIAL SOURCE	NVIDIA
Advanced Micro Devices (AMD) expects its revenue market sha...	vertexaisearch.cloud.google.com	OFFICIAL SOURCE	AMD Strategy Update
Lattice Semiconductor reported revenue of $201.08 million f...	vertexaisearch.cloud.google.com	OFFICIAL SOURCE	Lattice Semiconductor Corpo...
Total addressable market (TAM) for AI accelerators and serv...	vertexaisearch.cloud.google.com	OFFICIAL SOURCE	AMD / Futurum Research
AMD and Meta entered a multi-year partnership in February 2...	vertexaisearch.cloud.google.com	OFFICIAL SOURCE	AMD / Meta

DISCLAIMER

This research is based solely on the supplied Valuatum model, verified-gate materials and research register. Business-division revenue, profit, unit and valuation allocations are analyst estimates where AMD does not disclose product-level figures. The report is for institutional information, is not investment advice, and does not account for any reader’s objectives, constraints or tax position. Forecasts and target prices are uncertain and may change materially as operating, market or regulatory evidence develops.

This report contains an investment recommendation: a rating and a 12-month target price for the subject company's share. It is not personalised investment advice and takes no account of any individual recipient's objectives, financial situation or needs.

The target price is derived from Sum of the parts, EV/EBITDA, P/E, weighted as shown in the target price bridge, against peer-group and historical multiples. The valuation methodology, key assumptions and their sensitivities are presented in the valuation section of this report.

The report is updated on demand — on new company disclosures or at the distributing client's request — with no fixed update schedule. The report is produced as a commercial service; the distributing client may have a commercial relationship with Valuatum Oy.

This report was produced with the assistance of generative artificial intelligence. It may contain inaccuracies. Figures should be independently verified before use.

© Valuatum Oy